

Intermediate Chef

training@chef.io

Copyright (C) 2014 Chef Software, Inc.

v1.1.4

Introductions

v1.1.4

Introduce Yourselves

- Name
- Current job role
- Previous job roles/background
- Experience with Chef and/or config management
- Location
- What are you hoping to get out of the class?

Course Objectives and Style

v1.1.4

Course Objectives

- Upon completion of this course you will be able to:
 - Extend Chef with custom resources and providers
 - Understand the internals of a Chef Client run
 - Create, debug, and distribute custom Ohai plugins
 - Use the chef-shell debugger to interactively debug cookbooks
 - Configure report and exception handlers
 - Avoid common cookbook errors using Foodcritic and Rubocop
 - Gain an introduction to unit testing using ChefSpec

Course Pre-Requisites

- Completed Chef Fundamentals or equivalent experience
- You should install the most recent version of Chef Development Kit from here
 - <http://downloads.chef.io/chef-dk/>
- You can check your version with the command `chef -v`

How to learn Chef

- You bring the domain expertise about your business and problems
- Chef provides a framework for solving those problems
- Our job is to work together to teach you how to express solutions to your problems with Chef

Training is really a discussion

- We will be doing things the **hard way**
- We're going to do **a lot** of typing
- You can't be:
 - Absent
 - Late
 - Left Behind
- We will troubleshoot and fix bugs on the spot
- The result is you reaching fluency fast

Training is really a discussion

- I'll post objectives at the beginning of a section
- Ask questions when they come to you
- Ask for help when you need it
- You'll get the slides at the **end of class**

Agenda

v1.1.4

Topics

- Chef Fundamentals Refresher
- Building Custom Resources
- chef-shell - The Chef Debugger
- Writing Ohai Plugins
- Chef Client Run Internals
- Implementing Chef Handlers
- Correctness Checking Your Cookbooks
- An Introduction to ChefSpec
- Further Resources

Breaks!

- We'll take a break between each section, or every hour, whichever comes first
- We'll obviously break for lunch :)

Legend

v1.1.4

Legend: Do I run that command on my workstation?

This is an example of a command to run on your workstation

```
$ whoami  
i-am-a-workstation
```

This is an example of a command to run on your target node via SSH.

```
user@hostname:~$ whoami  
i-am-a-chef-node
```

Legend: Example Terminal Command and Output

```
$ ifconfig
```

```
lo0: flags=8049<UP,LOOPBACK,RUNNING,MULTICAST> mtu 16384
    options=3<RXCSUM,TXCSUM>
    inet6 fe80::1%lo0 prefixlen 64 scopeid 0x1
        inet 127.0.0.1 netmask 0xff000000
            inet6 ::1 prefixlen 128
gif0: flags=8010<POINTOPOINT,MULTICAST> mtu 1280
stf0: flags=0<> mtu 1280
en0: flags=8863<UP,BROADCAST,SMART,RUNNING,SIMPLEX,MULTICAST> mtu 1500
    ether 28:cf:e9:1f:79:a3
    inet6 fe80::2acf:e9ff:fe1f:79a3%en0 prefixlen 64 scopeid 0x4
        inet 10.100.0.84 netmask 0xffffffff broadcast 10.100.0.255
            media: autoselect
            status: active
p2p0: flags=8843<UP,BROADCAST,RUNNING,SIMPLEX,MULTICAST> mtu 2304
    ether 0a:cf:e9:1f:79:a3
    media: autoselect
    status: inactive
```

Legend: Example of editing a file on your workstation



OPEN IN EDITOR: ~/hello_world

Hi!

I am a friendly file.

SAVE FILE !

Training Environment (CloudShare)

v1.1.4

CloudShare Environment

- <https://use.cloudshare.com/Class/20022>
- Passphrase: **getcooking**

CloudShare Node

- Register and login to CloudShare (see invite)
- Start Using This Environment

Chef Intermediate Topics

Your dedicated hands-on environment is just a click away.

We believe you shouldn't have to waste time copying gigabytes of software, shipping machines, or traveling, just to get your IT into people's hands.

When you click the 'Start Using' button to your right, you'll have instant access worldwide to a full, enterprise-grade IT environment, available through your browser, mobile device, or dedicated client and VPN connection. You can connect your existing datacenter servers, change configurations, load new software, and provide this functionality to others. It's just like an on-premises data center, but with more flexibility and none of the headaches of backup, access, replication, logging, SLAs, and other annoyances - we take care of all of that for you.

CloudShare's SaaS platform is a quick and easy way to share copies of your complex IT environments, online. So you can collaborate with customers, partners, and colleagues - for Demos, Proofs-of-Concept, Training, or other enterprise applications.

Questions? Click [here](#) to email this environment's author or click [here](#) for CloudShare support.

Start Using
This Environment

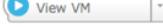


CloudShare Node

Chef Intermediate Topics

CentOS 6

CentOS 6.3 Server
Description: OS: CentOS 6.3 x64
Spec: 16 GB HD / 1 GB RAM
OS: Linux
State: Running
[More details >](#)

 [View VM](#)   

Chef Intermediate Topics

CentOS 6

CentOS 6.3 Server
Description: OS: CentOS 6.3 x64
Spec: 16 GB HD / 1 GB RAM
OS: Linux
State: Running
[Hide details ▾](#)

VM Details

External Address:	uv017z80u72k71dhkqt.vm.cld.sr	
Internal IP:	10.160.201.90	
Total Memory:	1024 MB	
Disk Size:	16 GB	
CPU:	1	

The machine was prepared in 2 minutes

Credentials (show password)

Auto-login:	root (local user)
Username:	root
Password:	*****

Lab - Login

```
$ ssh chef@<EXTERNAL_ADDRESS>
```

```
The authenticity of host 'uvolqrwls0jdgs3blvt.vm.cld.sr
(69.195.232.110)' can't be established.
RSA key fingerprint is d9:95:a3:b9:02:27:e9:cd:
74:e4:a2:34:23:f5:a6:8b.
Are you sure you want to continue connecting (yes/no)? yes
Warning: Permanently added 'uvolqrwls0jdgs3blvt.vm.cld.sr,
69.195.232.110' (RSA) to the list of known hosts.
chef@uvolqrwls0jdgs3blvt.vm.cld.sr's password:
Last login: Mon Jan  6 16:26:24 2014 from
host86-145-117-53.range86-145.btcentralplus.com
[chef@CentOS63 ~]$
```

Checkpoint

- At this point you should have
 - One virtual machine (VM) or server that you'll use for the lab exercises
 - The IP address or public hostname
 - An application for establishing an ssh connection
 - 'sudo' or 'root' permissions on the VM

Chef Fundamentals Refresher

v1.1.4



Lesson Objectives

- After this lesson the, you will be able to
 - Create a new Organization
 - Manually configure your certificates (.pem files) & knife.rb on your workstation so it can communicate with Chef Server
 - Bulk upload all cookbook, role, environment and data bag files to the Chef Server

Problem Statement

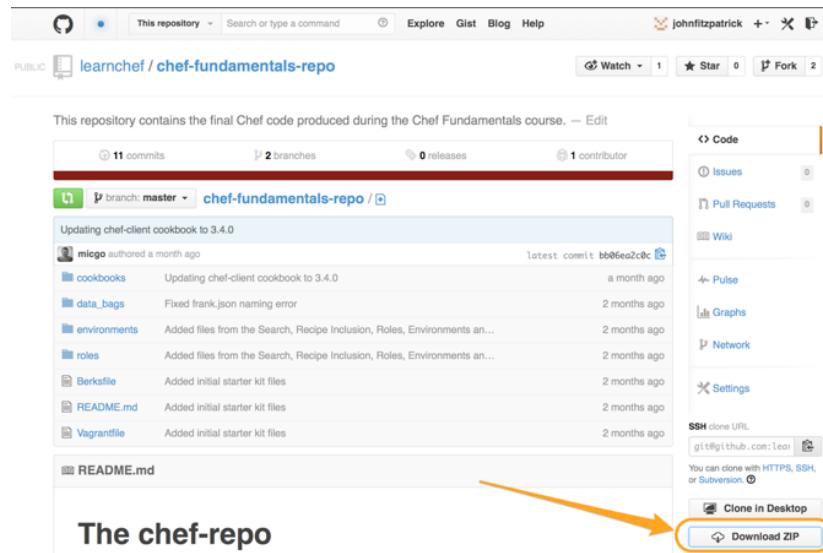
- **Problem:** We want to pick up where we left off after Chef Fundamentals class
- **Proposed Solution:** We need to
 - Set up a new Organization in Enterprise Chef
 - Configure your workstation to communicate with new Org
 - Download the Chef Fundamentals content from a GitHub repo
 - cookbooks, data bags, environments and roles
 - Upload all artifacts to your new Org
 - Configure server run list
 - Bootstrap server

Exercise: Set up a working directory

- For the purposes of this class make a working directory under your home directory called '~/intermediate', i.e.
 - Windows:-
 - C:\Users\you\intermediate
 - Mac/*nix:-
 - /Users/you/intermediate
- Navigate to this working Directory

Exercise: Download the Chef Fundamentals Repo

- Download the Chef Fundamentals working repo
<https://github.com/learnchef/chef-fundamentals-repo>



Exercise: Extract the repo to your working directory

```
$ cp ~/Downloads/chef-fundamentals-repo-master.zip .
$ unzip chef-fundamentals-repo-master.zip
```

```
Archive: chef-fundamentals-repo-master.zip
bb06ea2c0cabaa855e4cb1d1c43bbe4d75caf70d
  creating: chef-fundamentals-repo-master/
  inflating: chef-fundamentals-repo-master/Berksfile
  inflating: chef-fundamentals-repo-master/README.md
  inflating: chef-fundamentals-repo-master/Vagrantfile
  creating: chef-fundamentals-repo-master/cookbooks/
  creating: chef-fundamentals-repo-master/cookbooks/apache/
  inflating: chef-fundamentals-repo-master/cookbooks/apache/CHANGELOG.md
  inflating: chef-fundamentals-repo-master/cookbooks/apache/README.md
  creating: chef-fundamentals-repo-master/cookbooks/apache/attributes/
  creating: chef-fundamentals-repo-master
...
...
```

Exercise: View your working directory

```
$ cd chef-fundamentals-repo-master  
$ ls -a
```

```
.  Berksfile  Vagrantfile  data_bags      roles  
..  README.md  cookbooks   environments
```

- Notice there is no '.chef' directory here
- You need to create one and place your 'knife.rb' file, 'validator.pem' and your 'client.pem' in it

So what's in our working directory now?

```
cookbooks
├── apache
├── chef-client
├── cron
├── logrotate
├── motd
├── ntp
├── pci
└── starter
└── users

9 directories
```

```
environments
├── dev.rb
└── production.rb

2 files
```

```
roles
├── base.rb
├── starter.rb
└── webserver.rb

3 files
```

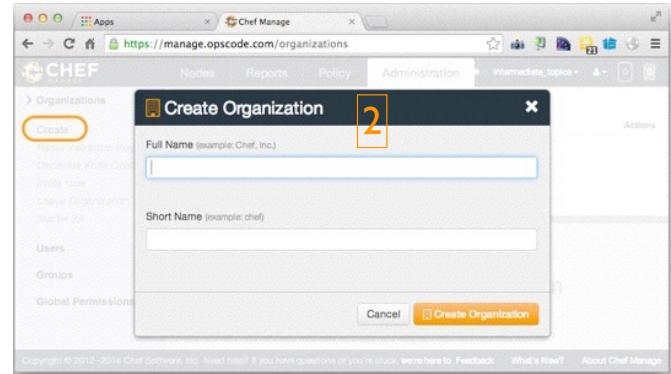
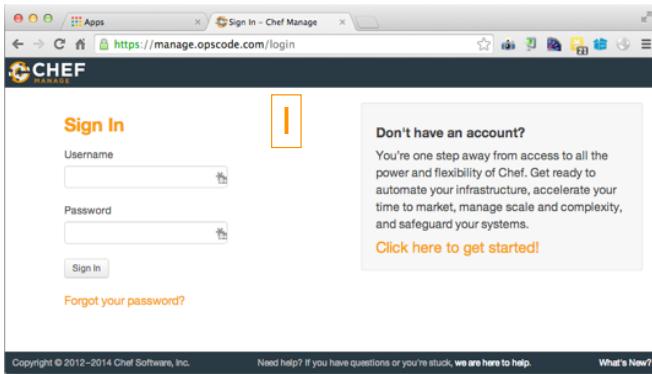
```
data_bags
├── groups
│   └── clowns.json
└── users
    ├── bobo.json
    └── frank.json

2 directories, 3 files
```

- These are the artifacts created in Chef Fundamentals
- Everything should be uploaded to the Chef Server

Exercise: Create a New Org

- Visit Hosted Enterprise Chef (manage.opscode.com)
- Sign in or create a new account
- Create a new Organization



Configuring your Workstation

- In Chef Fundamentals you set your workstation up the easy way using 'Starter Kit'
- In this class you will download your .pem files and knife.rb and configure it manually

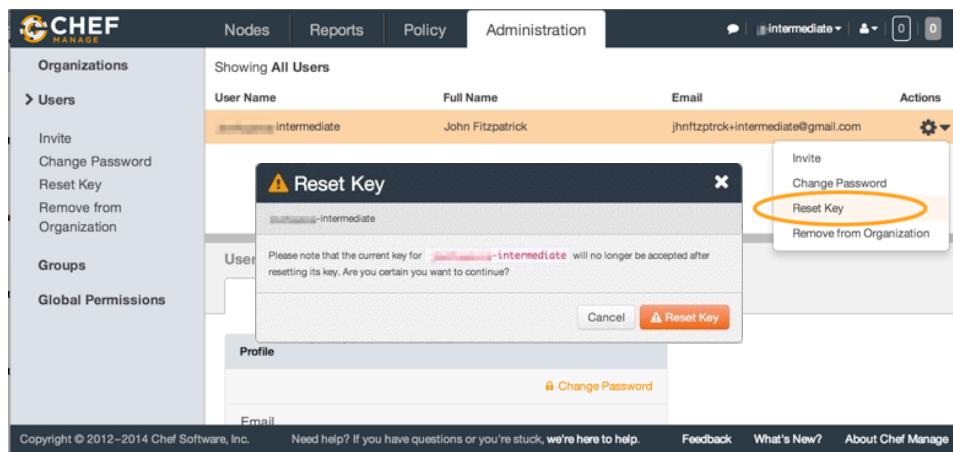
Exercise: Download knife.rb and validator pem

- Download the knife.rb and validator pem for your Org

The screenshot shows the Chef Manage interface. On the left, there's a sidebar with links like 'Organizations', 'Create', 'Reset Validation Key' (which is highlighted with an orange circle), 'Generate Knife Config' (also highlighted with an orange circle), 'Invite User', 'Leave Organization', and 'Starter Kit'. Below that are 'Users', 'Groups', and 'Global Permissions'. The main area shows 'Showing All Organizations' with one entry: 'jf-intermediate'. A modal window titled 'Reset Key' is open over this. The modal contains a large block of text representing a private key, starting with '-----BEGIN RSA PRIVATE KEY-----' and ending with '-----END RSA PRIVATE KEY-----'. At the bottom right of the modal are 'Close' and 'Download' buttons, with an orange arrow pointing to the 'Download' button.

Exercise: Download your client pem

- Reset and download your private client pem file
- **Only do this** if you don't already have the .pem file available on your laptop



Exercise: Create and populate a .chef directory

```
$ cd ~/intermediate/chef-fundamentals-repo-master  
$ mkdir .chef  
$ cp ~/Downloads/<yourname>.pem .chef  
$ cp ~/Downloads/<your-org>-validator.pem .chef  
$ cp ~/Downloads/knife.rb .chef
```

- knife.rb & .pem files reside in the .chef directory which can be in
 1. <current-directory>/ .chef
 2. /etc/.chef
 3. ~/ .chef

Exercise: Create and populate a .chef directory

What have we just done?

- You should now have these 3 files from Enterprise Chef
 - <yourname>.pem
 - <your-org>-validator.pem
 - knife.rb
- Later you will move these to your .chef directory

Exercise: Test your workstation

```
$ knife client list
```

```
<your-org>-validator.pem
```

Exercise: Upload Cookbooks

```
$ knife cookbook upload -a
```

```
Uploading apache          [ 0.2.0 ]
Uploading chef-client      [ 3.4.0 ]
Uploading cron             [ 1.2.8 ]
Uploading logrotate         [ 1.4.0 ]
Uploading motd              [ 0.1.0 ]
Uploading ntp                [ 1.5.4 ]
Uploading pci                [ 0.1.0 ]
Uploading starter            [ 1.0.0 ]
Uploading users              [ 0.1.0 ]
Uploaded all cookbooks.
```

Exercise: Upload data_bags

```
$ knife upload data_bags
```

```
Created data_bags/groups
Created data_bags/users
Created data_bags/groups/clowns.json
Created data_bags/users/bobo.json
Created data_bags/users/frank.json
```

Exercise: Upload Roles

```
$ knife role from file base.rb starter.rb webserver.rb
```

```
Updated Role base!
Updated Role starter!
Updated Role webserver!
```

Exercise: Upload Environments

```
$ knife environment from file dev.rb production.rb
```

```
Updated Environment dev
Updated Environment production
```

"Bootstrap" the Target Instance

```
$ knife bootstrap <EXTERNAL_ADDRESS> \
  --sudo -x chef -P chef -N "node1" --bootstrap-version 11.12.8
```

```
uv0164727i3mvh1jup2.vm.cld.sr --2014-05-13 04:31:10-- https://www.opscode.com/chef/install.sh
uv0164727i3mvh1jup2.vm.cld.sr Resolving www.opscode.com... 184.106.28.90
uv0164727i3mvh1jup2.vm.cld.sr Connecting to www.opscode.com|184.106.28.90|:443... connected.
uv0164727i3mvh1jup2.vm.cld.sr HTTP request sent, awaiting response... 200 OK
uv0164727i3mvh1jup2.vm.cld.sr Length: 15934 (16K) [application/x-sh]
uv0164727i3mvh1jup2.vm.cld.sr Saving to: `STDOUT'
uv0164727i3mvh1jup2.vm.cld.sr
100%[=====] 15,934      --.-K/s   in 0s
uv0164727i3mvh1jup2.vm.cld.sr
uv0164727i3mvh1jup2.vm.cld.sr 2014-05-13 04:31:10 (538 MB/s) - written to stdout [15934/15934]
uv0164727i3mvh1jup2.vm.cld.sr
uv0164727i3mvh1jup2.vm.cld.sr Downloading Chef 11.8.2 for el...
uv0164727i3mvh1jup2.vm.cld.sr downloading https://www.opscode.com/chef/metadata?
v=11.8.2&prerelease=false&p=el&pv=6&m=x86_64
uv0164727i3mvh1jup2.vm.cld.sr to file /tmp/install.sh.41533/metadata.txt
uv0164727i3mvh1jup2.vm.cld.sr trying wget...
uv0164727i3mvh1jup2.vm.cld.sr url https://opscode-omnibus-packages.s3.amazonaws.com/el/6/x86_64/
chef-11.8.2-1.el6.x86_64.rpm
...
```

knife bootstrap



Workstation



Node

knife bootstrap

```
knife bootstrap HOSTNAME -x root -P PASSWORD -
```



Chef Server



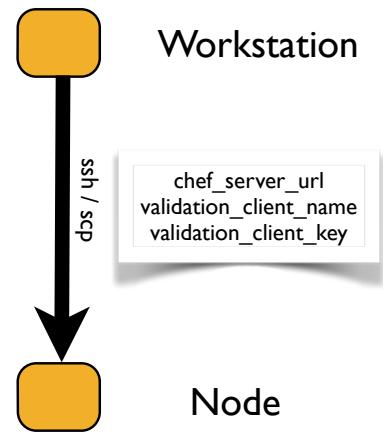
Workstation



Node

knife bootstrap

```
knife bootstrap HOSTNAME -x root -P PASSWORD -
```



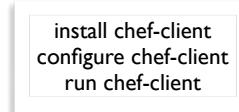
knife bootstrap

```
knife bootstrap HOSTNAME -x root -P PASSWORD -
```



Workstation

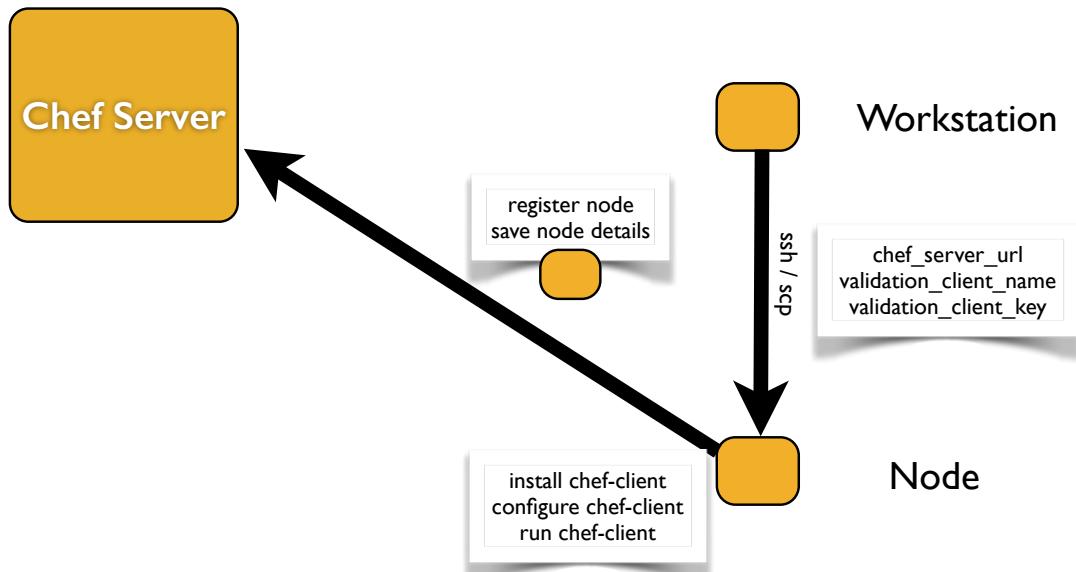
chef_server_url
validation_client_name
validation_client_key



Node

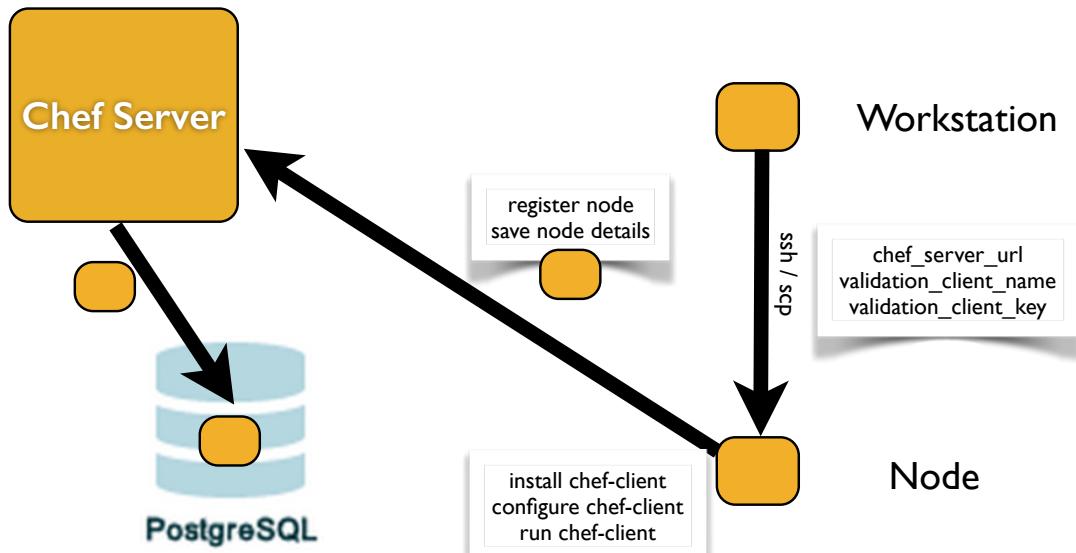
knife bootstrap

```
knife bootstrap HOSTNAME -x root -P PASSWORD -
```



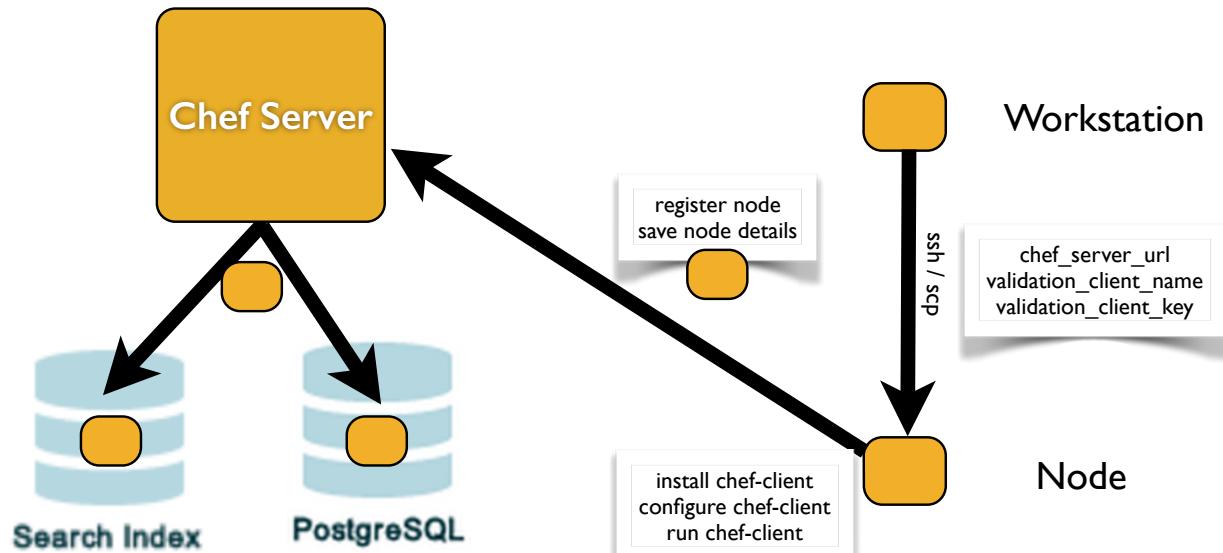
knife bootstrap

```
knife bootstrap HOSTNAME -x root -P PASSWORD -
```



knife bootstrap

```
knife bootstrap HOSTNAME -x root -P PASSWORD -
```



What just happened?

- Chef and all of its dependencies installed via an operating system-specific package ("omnibus installer")
- Installation includes
 - The Ruby language - used by Chef
 - knife - Command line tool for administrators
 - chef-client - Client application
 - ohai - System profiler
 - ...and more

Verify Your Target Instance's Chef-Client is Configured Properly

```
$ ssh chef@<EXTERNAL_ADDRESS>

chef@node1:~$ ls /etc/chef
client.pem  client.rb  first-boot.json validation.pem

chef@node1:~$ which chef-client
/usr/bin/chef-client

chef@node1:~$ chef-client -v
Chef: 11.12.8
```

Exercise: Set run list for the managed node

```
$ knife node run_list set node1 'role[webserver]'
```

- For Powershell you need to do this

```
$ knife node run_list set node1 `'role[webserver]'`
```

```
node1:  
  run_list:  
    role[webserver]
```

Exercise: Re-run the Chef Client

```
chef@node1$ sudo chef-client -l info
```

```
...
[2014-01-07T06:57:48-05:00] INFO: template[/etc/httpd/conf.d/bears.conf] updated file contents /etc/
httpd/conf.d/bears.conf

- update content in file /etc/httpd/conf.d/bears.conf from 30610b to 92e1e4
--- /etc/httpd/conf.d/bears.conf2014-01-07 05:13:42.450113970 -0500
+++ /tmp/chef-rendered-template20140107-18263-5563t5 2014-01-07 06:57:48.091397610 -0500
@@ -1,6 +1,6 @@
- Listen 81
+ Listen 8081

-<VirtualHost *:81>
+<VirtualHost *:8081>
    ServerAdmin webmaster@localhost

    DocumentRoot /srv/apache/bears
...

```

SSL Problem?

- You may have seen an SSL error

```
* * * * * * * * * * * * * * * * * * * * * * * * * * * * * * * * * * * * * * * * * *  
SSL validation of HTTPS requests is disabled. HTTPS connections are still  
encrypted, but chef is not able to detect forged replies or man in the middle  
attacks.
```

To fix this issue add an entry like this to your configuration file:

```
```  
Verify all HTTPS connections (recommended)
ssl_verify_mode :verify_peer

OR, Verify only connections to chef-server
verify_api_cert true
```
```

To check your SSL configuration, or troubleshoot errors, you can use the
`knife ssl check` command like so:

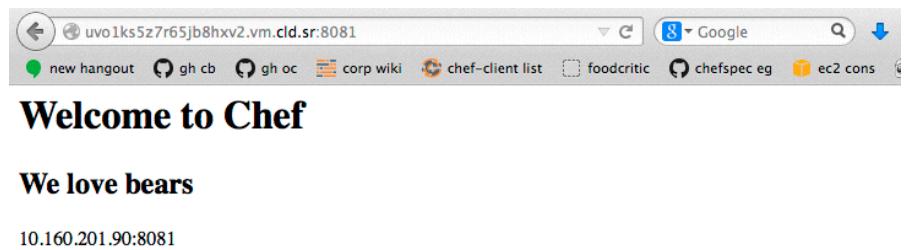
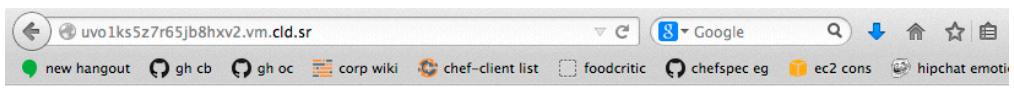
- This is harmless for the purpose of this course
- But, we will fix it in the next section!

Exercise: Verify chef-client is running

```
chef@node1$ ps awux | grep chef-client
```

```
root      8933  0.3  2.2 130400 37816 ?          S1
03:19  0:01 /opt/chef/embedded/bin/ruby /usr/bin/
chef-client -d -c /etc/chef/client.rb -L /var/log/
chef/client.log -P /var/run/chef/client.pid -i 1800
-s 300
```

Exercise: Verify the two sites are working!



Review Questions

- What files did we download from Chef Server to configure our workstation?
- Where did we place these files?
- What other way could we have copied the repo files from GitHub?
 - Bonus point - why did we not do this?
- How did we upload all cookbooks to the Chef Server?

Configuring Chef Clients

SSL Verification and chef-client management

v1.1.4

Lesson Objectives

- After completing the lesson, you will be able to
 - Use the chef-client cookbook to manage the Chef Client service on your nodes
 - Configure nodes to use SSL Certificate Validation

The Problem and Success Criteria

- **The Problem:** We want to make sure the Chef Client service is started and enabled on boot. We also want to make sure Chef Client is validating all SSL connections.
- **Success Criteria:** We will use the `chef-client` cookbook to accomplish these tasks.

Using the chef-client cookbook

- The chef-client cookbook manages the Chef Client's:
 - `client.rb` configuration
 - service (runit, systemd, init, SMF, etc.)
 - validation key deletion

Examining the chef-client cookbook

- We're already using two recipes on the node from the chef-client cookbook
 - `chef-client::service`
(via `chef-client::default`)
 - `chef-client::delete_validation`

Exercise: View node1's Run List

```
$ knife node show node1
```

```
Node Name:    node1
Environment:  _default
FQDN:        centos63.example.com
IP:          10.160.201.90
Run List:    role[webserver]
Roles:       webserver, base
Recipes:     chef-client::delete_validation, chef-client, ntp, motd, users,
             apache, chef-client::default, chef-client::service, chef-client::init_service,
             ntp::default, motd::default, users::default, users::groups, apache::default
Platform:    centos 6.4
Tags:
```

Exercise: View the chef-client::config recipe



OPEN IN EDITOR: cookbooks/chef-client/recipes/config.rb

```
template "#{node['chef_client']['conf_dir']}/client.rb" do
  source 'client.rb.erb'
  owner d_owner
  group d_group
  mode 00644
  variables(
    :chef_config => node['chef_client']['config'],
    :chef_requires => chef_requires,
    :ohai_disabled_plugins => node['ohai']['disabled_plugins'],
    :start_handlers => node['chef_client']['config']['start_handlers'],
    :report_handlers => node['chef_client']['config']['report_handlers'],
    :exception_handlers => node['chef_client']['config']
  ['exception_handlers'])
  )
  notifies :create, 'ruby_block[reload_client_config]', :immediately
end
```

Managing client.rb settings

- The config recipe is used to dynamically generate the `/etc/chef/client.rb` config file
- The template walks all attributes in `node['chef_client']['config']` and writes them out as key:value pairs
- The key should be the configuration directive
- The `chef_server_url`, `node_name` and `validation_client_name` are set by default in the attributes file from `Chef::Config`

Managing client.rb settings

- For example, the following attributes (in a role):

```
default_attributes(  
  "chef_client" => {  
    "config" => {  
      "ssl_verify_mode" => ":verify_peer",  
      "log_level" => ":info"  
    }  
  }  
)
```

- Will render the following configuration (/etc/chef/client.rb):

```
chef_server_url "https://api.opscode.com/organizations/MYORG"  
validation_client_name "MYORG-validator"  
ssl_verify_mode :verify_peer  
node_name "config-ubuntu-1204"  
log_level :info
```

SSL Certificate Validation

- Chef defaults to not verifying certificates when it makes HTTPS connections - hence the error when running chef-client
- Enable SSL Verification by setting `ssl_verify_mode :verify_peer` in your config file
- Two useful knife commands:-
 - **knife ssl check** - makes an SSL connection to your Chef server or any other HTTPS server and tells you if the server presents a valid certificate
 - **knife ssl fetch** - allows you to automatically fetch a server's certificates to your trusted certs directory

Log Levels

- chef-client defaults to the log level :**auto**
- chef-client supports additional levels
 - :**debug**
 - :**info**
 - :**warn**
 - :**error**
 - :**fatal**
- Reference: https://docs.chef.io/config_rb_client.html

Exercise: Add chef-client::config to the base role



OPEN IN EDITOR: roles/base.rb

```
name "base"
description "Base Server Role"
run_list "recipe[chef-client::delete_validation]",
"recipe[chef-client::config]", "recipe[chef-client]",
"recipe[ntp]", "recipe[motd]", "recipe[users]"
default_attributes({
  "chef_client" => {
    "config" => {
      "ssl_verify_mode" => ":verify_peer",
      "log_level" => ":info"
    }
  }
})
```

SAVE FILE!

Exercise: Upload the base role

```
$ knife role from file base.rb
```

Updated Role base!

Exercise: Run chef-client

```
chef@node1$ sudo chef-client
```

```
Starting Chef Client, version 11.12.8
resolving cookbooks for run list: ["chef-client::delete_validation", "chef-client::config",
"chef-client", "ntp", "motd", "users", "apache"]
Synchronizing Cookbooks:
- chef-client
- cron
- logrotate
- ntp
- motd
- pci
- users
- apache
- ohai
Compiling Cookbooks...
...
```

Exercise: Verify config updated

```
chef@node1$ sudo cat /etc/chef/client.rb
```

```
...
ssl_verify_mode :verify_peer
log_level :info
...
```

Special Handling for Non-CA Signed Certificates

- Hosted Chef's certificate is signed by a known CA
- However, if you run your own CA, or your Chef Server's certificate is self-signed, you can use knife to import it

Exercise: Fetch Hosted Chef's SSL certificate

```
$ knife ssl fetch
```

WARNING: Certificates from api.opscode.com will be fetched and placed in your trusted_cert directory (C:/Users/you/.chef/trusted_certs).

Knife has no means to verify these are the correct certificates. You should verify the authenticity of these certificates after downloading.

```
Adding certificate for *.opscode.com in C:/Users/you/.chef/trusted_certs/wildcard_opscode_com.crt
```

```
Adding certificate for DigiCert Secure Server CA in C:/Users/you/.chef/trusted_certs/DigiCert_Secure_Server_CA.crt
```

Review Questions

- How do you remove the validation .pem on the node?
- What recipe is used to configure chef-client?
- How do you enable SSL Verification on chef-client?
- How do you fetch a Chef Server's certificates to your trusted certs directory?

Building Custom Resources

Expanding functionality in the Chef Framework

v1.1.4

Lesson Objectives

- After completing the lesson, you will be able to
 - Describe the components of the Lightweight Resource/Provider (LWRP) framework
 - Explain the Resource DSL
 - Explain the Provider DSL
 - Build a resource and provider from scratch
 - Look at examples of other LWRPs

A Brief Review...

- **Resources** are declarative interfaces that describe *what* we want to happen, rather than *how*
- Resources take action through **Providers** that perform the *how*
- Resources
 - have a **type**
 - have a **name**
 - can have one or more **parameters**
 - take **action** to put the resource into a desired state
 - can send **notifications** to other resources

Other Ways To Extend Chef

- **Definitions**

- “recipe macro”
- stored in definitions/
- cannot receive notifications

- **Heavyweight Resources**

- pure Ruby code
- stored in libraries/
- cannot use core resources (by default)

Components of an LWRP

- LWRPs have two components: the resource and the provider
 - **Resources** are used in Recipes to declare the state to configure on our system
 - **Providers** configure that state on the system during convergence
- They are defined in the `resources/` and `providers/` directories of cookbooks

The Problem and the Success Criteria

- **The Problem:** We want to abstract the Apache virtual host configuration pattern.
- **Success Criteria:** We will create an `apache_vhost` resource that will let us manage Apache virtual hosts.

Exercise: Change the cookbook's version number in the metadata



OPEN IN EDITOR: cookbooks/apache/metadata.rb

```
maintainer      "apache"
maintainer_email "YOUR_EMAIL"
license         "All rights reserved"
description     "Installs/Configures apache"
long_description IO.read(File.join(File.dirname(__FILE__), 'README.md'))
version         "0.3.0"
```

SAVE FILE!

- Major, Minor, Patch
- Semantic Versioning Policy: <http://semver.org/>

Resource & Provider Naming

```
+ cookbooks
  |
  + apache
    |
    + resources
      |
      + vhost.rb
```

- Name of new resource is implied by its name in the cookbook
- In the above scenario, the new resource will be called `apache_vhost`

The Resource DSL

- Three methods: `actions`, `attribute`, `default_action`
 - `actions` defines a list of allowed actions by the resource
 - `attribute` defines a new parameter for the resource block
 - `default_action` defines the one action to use if no action is specified in the resource block

Exercise: Review the apache::default recipe



OPEN IN EDITOR: cookbooks/apache/recipes/default.rb

```
# Iterate over the apache sites
node["apache"]["sites"].each do |site_name, site_data|
  # Set the document root
  document_root = "/srv/apache/#{site_name}"
```

- Note the file cookbooks/apache/attributes/default.rb contains the following

```
default["apache"]["sites"]["clowns"] = { "port" => 80 }
default["apache"]["sites"]["bears"] = { "port" => 81 }
```

Exercise: Review the apache::default recipe



OPEN IN EDITOR: cookbooks/apache/recipes/default.rb

```
# Add a template for Apache virtual host configuration
template "/etc/httpd/conf.d/#{site_name}.conf" do
  source "custom.erb"
  mode "0644"
  variables(
    :document_root => document_root,
    :port => site_data["port"]
  )
  notifies :restart, "service[httpd]"
end
```

Exercise: Review the apache::default recipe



OPEN IN EDITOR: cookbooks/apache/recipes/default.rb

```
# Add a directory resource to create the document_root
directory document_root do
  mode "0755"
  recursive true
end
```

Exercise: Review the apache::default recipe



OPEN IN EDITOR: cookbooks/apache/recipes/default.rb

```
# Add a template resource for the virtual host's
index.html
template "#{document_root}/index.html" do
  source "index.html.erb"
  mode "0644"
  variables(
    :site_name => site_name,
    :port => site_data["port"]
  )
end
```

Exercise: Create an apache_vhost resource with two allowed actions

 **OPEN IN EDITOR:** cookbooks/apache/resources/vhost.rb

```
actions :create, :remove
```

SAVE FILE !

- We have allowed two resource actions

The Provider DSL

- One method: `action`
 - Specify the `action` name with a Ruby symbol
 - Name maps to allowed actions defined in the resource actions
- You can also re-use any Chef Resources inside your providers

Exercise: Create provider for the :create action



OPEN IN EDITOR: cookbooks/apache/providers/vhost.rb

```
action :create do
  puts "My name is #{new_resource.name}"
end
```

SAVE FILE!

- When our apache_vhost resource calls action :create, execute this block.

Exercise: Set an action in our apache::default recipe



OPEN IN EDITOR: cookbooks/apache/recipes/default.rb

```
# Disable the default virtual host
execute "mv /etc/httpd/conf.d/
welcome.conf /etc/httpd/conf.d/
welcome.conf.disabled" do
  only_if do
    File.exist?("/etc/httpd/conf.d/
welcome.conf")
  end
  notifies :restart, "service[httpd]"
end

# Enable an Apache Virtualhost
apache_vhost "lions" do
  action :create
end
```

- We call resource **apache_vhost**
- With name "**lions**"
- With action **:create** in the parameter block

SAVE FILE !

Exercise: Upload the apache cookbook

```
$ knife cookbook upload apache
```

```
Uploading apache      [ 0.3.0 ]
Uploaded 1 cookbook.
```

Exercise: Run chef-client

```
chef@node1$ sudo chef-client
```

```
Starting Chef Client, version 11.10.4
resolving cookbooks for run list: ["chef-client::delete_validation", "chef-client", "ntp", "motd", "users", "apache"]
Synchronizing Cookbooks:
  - chef-client
  - cron
  - logrotate
  - ntp
  - motd
  - pci
  - users
  - apache
...
  * apache_vhost[lions] action create[My name is lions
  (up to date)
...
Running handlers:
Running handlers complete

Chef Client finished, 12/35 resources updated in 12.045321863 seconds
```

Exercise: Use a Chef Resource within your provider



OPEN IN EDITOR: cookbooks/apache/providers/vhost.rb

```
use_inline_resources

action :create do
  log "My name is #{new_resource.name}"
end
```

SAVE FILE!

- The `log` resource uses Chef's logger object to print messages at `Chef::Config[:log_level]`

use_inline_resources

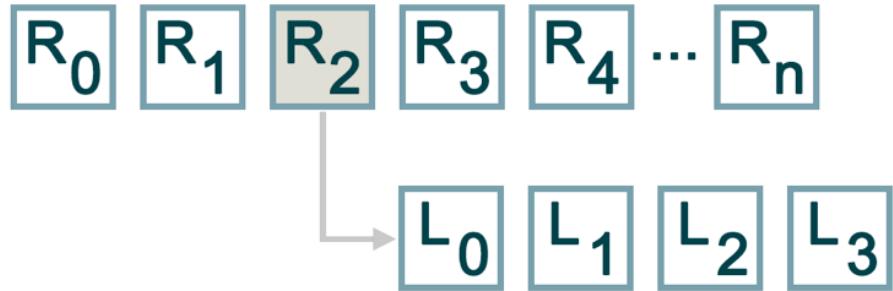
- `use_inline_resources` instructs the embedded resources (“`log`”) to notify the parent (“`apache_vhost`”) if their states change
- More info: [**docs.chef.io/
lwrp_custom_provider.html#use-inline-resources**](https://docs.chef.io/lwrp_custom_provider.html#use-inline-resources)

LWRPs and the Resource Collection

R_0 R_1 R_2 R_3 R_4 ... R_n

- A regular resource collection is just a big array of resources.

LWRPs and the Resource Collection



- `use_inline_resources` creates a mini resource collection (execution context)
- Notifications are rolled up to the master resource collection.

Exercise: Upload the apache cookbook

```
$ knife cookbook upload apache
```

```
Uploading apache [ 0.3.0 ]
Uploaded 1 cookbook.
```

Exercise: Run chef-client

```
chef@node1$ sudo chef-client
```

```
Starting Chef Client, version 11.10.4
resolving cookbooks for run list: ["chef-client::delete_validation", "chef-client", "ntp", "motd", "users", "apache"]
Synchronizing Cookbooks:
  - chef-client
  - cron
  - logrotate
  - ntp
  - motd
  - pci
  - users
  - apache
...
* apache_vhost[lions] action createRecipe: <Dynamically Defined Resource>
* log[My name is lions] action write
...
Running handlers:
Running handlers complete

Chef Client finished, 12/35 resources updated in 12.045321863 seconds
```

Exercise: Create attribute parameters for the apache_vhost resource



OPEN IN EDITOR: cookbooks/apache/resources/vhost.rb

```
actions :create, :remove

attribute :site_name, :name_attribute => true, :kind_of => String
attribute :site_port, :default => 80, :kind_of => Fixnum
```

SAVE FILE!

- attribute takes the name of the attribute and an (optional) hash of validation parameters
- These validation parameters are specific to the LWRP Resource DSL

Resource Validation Parameters



OPEN IN EDITOR: cookbooks/apache/resources/vhost.rb

```
actions :create, :remove

attribute :site_name, :name_attribute => true, :kind_of => String
attribute :site_port, :default => 80, :kind_of => Fixnum
```

SAVE FILE!

- **:name_attribute** set the site_name to the resource name
- **:default** sets the default value for this parameter
- **:kind_of** ensures the value is a particular class

Resource Validation Parameters

Validation	Meaning
:callbacks	Hash of Procs, should return true
:default	Default value for the parameter
:equal_to	Match the value with ==
:kind_of	Ensure the value is of a particular class
:name_attribute	Set to the resource name
:regex	Match the value against the regex
:required	The parameter must be specified
:respond_to	Ensure the value has a given method

:kind_of examples

- `:kind_of` accepts many types, either built-in Ruby classes, or your own

- - `:kind_of => String # String`
 - `:kind_of => Array # Array`
 - `:kind_of => Fixnum # Fixnum`
 - `:kind_of => [:some, :list, :of, :symbols]`
 - `:kind_of => [TrueClass, FalseClass]`
 - `:kind_of => [String, Array] # composite`

Exercise: Extend :create action



OPEN IN EDITOR: cookbooks/apache/providers/vhost.rb

```
use_inline_resources

action :create do
  # Set the document root
  document_root = "/srv/apache/#{new_resource.site_name}"

  # Add a template for Apache virtual host configuration
  template "/etc/httpd/conf.d/#{new_resource.site_name}.conf" do
    source "custom.erb"
    mode "0644"
    variables(
      :document_root => document_root,
      :port => new_resource.site_port
    )
  end
end
```

SAVE FILE !

106

Exercise: Extend :create action



OPEN IN EDITOR: cookbooks/apache/providers/vhost.rb

```
# Add a directory resource to create the document_root
directory document_root do
  mode "0755"
  recursive true
end

# Add a template resource for the virtual host's index.html
template "#{document_root}/index.html" do
  source "index.html.erb"
  mode "0644"
  variables(
    :site_name => new_resource.site_name,
    :port => new_resource.site_port
  )
end
end
```

SAVE FILE !

107

Exercise: Use apache_vhost in a recipe



OPEN IN EDITOR: cookbooks/apache/recipes/default.rb

```
# Enable an Apache Virtualhost
apache_vhost "lions" do
  site_port 8080
  action :create
  notifies :restart, "service[httpd]"
end

# Iterate over the apache sites
node["apache"]["sites"].each do |site_name, site_data|
```

SAVE FILE !

Exercise: Upload the apache cookbook

```
$ knife cookbook upload apache
```

```
Uploading apache [ 0.3.0 ]  
Uploaded 1 cookbook.
```

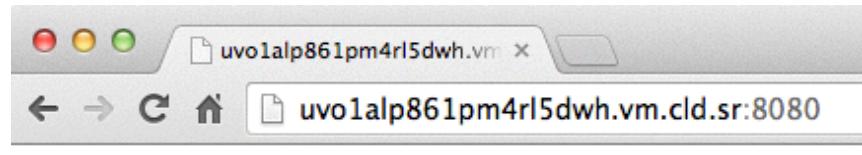
Exercise: Run chef-client

```
chef@node1$ sudo chef-client
```

```
Starting Chef Client, version 11.10.4
resolving cookbooks for run list: ["chef-client::delete_validation", "chef-client", "ntp", "motd", "users", "apache"]
Synchronizing Cookbooks:
  - chef-client
  - cron
  - logrotate
  - ntp
  - motd
  - pci
  - users
  - apache
...
* apache_vhost[lions] action createRecipe: <Dynamically Defined Resource>
* template[/etc/httpd/conf.d/lions.conf] action create
...
Running handlers:
Running handlers complete

Chef Client finished, 7/38 resources updated in 9.724844315 seconds
```

Exercise: Verify new ‘lions’ site



We love lions

10.160.157.59:8080

The resource collection - inline resources

Resource Collection

```
resource_collection = [
    ...,
    package[httpd],
    service[httpd],
    apache_vhost[lions],
    resource_collection = [
        template["/etc/httpd/conf.d/lions.conf"],
        directory["/srv/apache/lions"],
        template["/srv/apache/lions/index.html"],
    ],
    template["/etc/httpd/conf.d/clowns.conf"],
    directory["/srv/apache/clowns"],
    template["/srv/apache/clowns/index.html"],
    template["/etc/httpd/conf.d/bears.conf"],
    directory["/srv/apache/bears"],
    template["/srv/apache/clowns/bears.html"]
]
```

All notifications in inline resources are rolled up to 'master' resource collection

Houston, we have a problem!

- Now that we have an easy way to add virtual hosts, we will quickly run into a data management problem
- We could keep adding more attributes to the apache cookbook, but...
 - that doesn't scale well across teams
 - you can end up with a long, messy list of attributes

Let's use Data Bags to drive our new LWRP

```
default[ "apache" ][ "sites" ][ "clowns" ] = { "port" => 80 }
default[ "apache" ][ "sites" ][ "bears" ] = { "port" => 81 }
```

- How should we structure our data bags?
 - Data bag for each site?
 - One data bag for apache with items for each site?

Exercise: Create the apache_sites data bag

```
$ knife data_bag create apache_sites
```

```
Created data_bag[apache_sites]
```

Exercise: Create the apache_sites data bag

```
$ mkdir -p data_bags/apache_sites
```

No output...

Exercise: Create clowns data bag item



OPEN IN EDITOR: data_bags/apache_sites/clowns.json

```
{  
  "id": "clowns",  
  "port": 80  
}
```

SAVE FILE !

117

Exercise: Upload the clowns data bag item

```
$ knife data_bag from file apache_sites clowns.json
```

```
Updated data_bag_item[apache_sites::clowns]
```

Exercise: Create bears data bag item

 **OPEN IN EDITOR:** data_bags/apache_sites/bears.json

```
{  
  "id": "bears",  
  "port": 8081  
}
```

SAVE FILE !

119

Exercise: Upload the bears data bag item

```
$ knife data_bag from file apache_sites bears.json
```

```
Updated data_bag_item[apache_sites::bears]
```

Exercise: Create lions data bag item



OPEN IN EDITOR: data_bags/apache_sites/lions.json

```
{  
  "id": "lions",  
  "port": 8080  
}
```

SAVE FILE !

121

Exercise: Upload the bears data bag item

```
$ knife data_bag from file apache_sites lions.json
```

```
Updated data_bag_item[apache_sites::lions]
```

Exercise: Refactor apache::default recipe



OPEN IN EDITOR: cookbooks/apache/recipes/default.rb

```
# Iterate over the apache sites
all_sites = search("apache_sites", "*:*")
all_sites.each do |site|
  # Enable an Apache Virtualhost
  apache_vhost site['id'] do
    site_port site['port']
    action :create
    notifies :restart, "service[httpd]"
  end
end
```

- Delete existing node[“apache”][“sites”] loop
- Move apache_vhost LWRP inside a search loop
- Change variable names

Exercise: Upload the apache cookbook

```
$ knife cookbook upload apache
```

```
Uploading apache [ 0.3.0 ]
Uploaded 1 cookbook.
```

Exercise: Run chef-client

```
chef@node1$ sudo chef-client
```

```
Starting Chef Client, version 11.10.4
resolving cookbooks for run list: ["chef-client::delete_validation", "chef-client", "ntp", "motd", "users", "apache"]
Synchronizing Cookbooks:
  - chef-client
  - cron
  - logrotate
  - ntp
  - motd
  - pci
  - users
  - apache
...
* apache_vhost[clowns] action createRecipe: <Dynamically Defined Resource>
..
* apache_vhost[bears] action createRecipe: <Dynamically Defined Resource>
..
* apache_vhost[lions] action createRecipe: <Dynamically Defined Resource>
..
Chef Client finished, 2/39 resources updated in 7.004059993 seconds
```

Exercise: Extend :remove action



OPEN IN EDITOR: cookbooks/apache/providers/vhost.rb

```
# Add a template resource for the virtual host's index.html
template "#{document_root}/index.html" do
  source "index.html.erb"
  mode "0644"
  variables(
    :site_name => new_resource.site_name,
    :port => new_resource.site_port
  )
end
end

action :remove do
  file "/etc/httpd/conf.d/#{new_resource.site_name}.conf" do
    action :delete
  end
end
```

SAVE FILE !

126

Exercise: Refactor apache::default recipe



OPEN IN EDITOR: cookbooks/apache/recipes/default.rb

```
# Disable the default virtual host
apache_vhost "welcome" do
  action :remove
  notifies :restart, "service[httpd]"
end

# Iterate over the apache sites
search("apache_sites", "*:*").each do |site|
  # Enable an Apache Virtualhost
  apache_vhost site['id'] do
```

- Delete existing **execute** resource that disables the welcome site
- Add a new **apache_vhost** resource with **action :remove**

Exercise: Upload the apache cookbook

```
$ knife cookbook upload apache
```

```
Uploading apache [ 0.3.0 ]
Uploaded 1 cookbook.
```

Exercise: Run chef-client

```
chef@node1$ sudo chef-client
```

```
Starting Chef Client, version 11.10.4
resolving cookbooks for run list: ["chef-client::delete_validation", "chef-client", "ntp", "motd", "users",
"apache"]
Synchronizing Cookbooks:
- chef-client
- cron
- logrotate
- ntp
- motd
- pci
- users
- apache
..
* apache_vhost[welcome] action removeRecipe: <Dynamically Defined Resource>
* file[/etc/httpd/conf.d/welcome.conf] action delete (up to date)
(up to date)
..
```

```
Chef Client finished, 2/39 resources updated in 7.004059993 seconds
```

Other ways to write resources & providers

- We wrote an LWRP using out-of-the-box Chef resources
- You can write pure Ruby in LWRPs as well
 - http://docs.chef.io/lwrp_custom_provider_ruby.html
- You can also write pure Chef resources/providers by inheriting from `Chef::Resource::Base` and `Chef::Provider::Base` and putting them in libraries/

Use Cases for LWRPs

- Hide unnecessary implementation details
- Example: app admin team decides how to create Apache virtual hosts, JBoss JVMs, etc.
- Write LWRPs to allow people to safely instantiate them
 - Parameter validation for correctness/conformance (e.g. “we don’t name JVMs with punctuation characters”)
 - Reject invalid values (e.g. “port -1”)
 - Restrict what tunables are available to consumers

Review Questions

- What are the two components of an LWRP?
- What is the difference between these two components?
- What language can LWRPs written in?
- What classes can you inherit from Chef to write HWRPs?

chef-shell - The Chef Debugger

A Read-Eval-Print-Loop (REPL) for Chef

v1.1.4

133



Lesson Objectives

- After completing the lesson, you will be able to
 - Use chef-shell, the interactive Chef console (REPL)
 - List chef-shell's execution contexts
 - Use chef-shell to inspect and debug recipes
 - Interact with a Chef Server API
 - Discuss the magic that is knife exec

Chef Shell

- Interactive Ruby console (REPL) for Chef
- Three startup modes (standalone, solo, client)
- Three run-time contexts (main, attribute, recipe)
- Configured with a specific or implied config file

Chef Shell Startup

- Three modes for startup
 - **Standalone:** no cookbooks and empty run list (default)
 - **Solo:** attempts to act as chef-solo, using solo config and JSON attributes. (-s or --solo)
 - **Client:** contacts chef server for node object, retrieving cookbooks and using specified run list (-z or --client).

Exercise: Start the chef-shell REPL

```
chef@node1$ sudo chef-shell -z
```

```
loading configuration: /etc/chef/client.rb
Session type: client
Loading.....resolving cookbooks for run list: ["apache::ohai_plugin", "chef-client::delete_validation", "chef-client", "ntp", "motd", "users",
"apache"]
..Synchronizing Cookbooks:
- apache
- chef-client
- cron
- logrotate
- ntp
- motd
- pci
- users
./opt/chef/embedded/lib/ruby/gems/1.9.1/gems/ohai-7.0.4/lib/ohai/plugins/linux/network.rb:49: warning: already initialized constant
IPROUTE_INT_REGEX
.....[2014-05-14T06:27:24-04:00] WARN: Cloning resource attributes for apache_vhost[lions] from prior resource (CHEF-3694)
[2014-05-14T06:27:24-04:00] WARN: Previous apache_vhost[lions]: /var/chef/cache/cookbooks/apache/recipes/default.rb:27:in `from_file'
[2014-05-14T06:27:24-04:00] WARN: Current apache_vhost[lions]: /var/chef/cache/cookbooks/apache/recipes/default.rb:42:in `block in from_file'
done.

This is the chef-shell.
Chef Version: 11.12.4
http://www.opscode.com/chef
http://docs.chef.io/
```

Exercise: Get help

```
chef > help
```

```
chef-shell Help
=====
| Command           | Description
=====
| help              | prints this help message
| version            | prints information about chef
| recipe_mode        | switch to recipe mode
| attributes_mode    | switch to attributes mode
| run_chef           | run chef using the current recipe
| chef_run            | returns an object to control a paused chef run
| chef_run.resume     | resume the chef run
| chef_run.step       | run only the next resource
| chef_run.skip_back  | move back in the run list
| chef_run.skip_forward | move forward in the run list
| reset              | resets the current recipe
```

Exercise: Get help

```
chef > help(:nodes)
```

```
Command: nodes
=====
## SUMMARY ##
+nodes+ Allows you to query your chef server for information about your nodes.

## LIST ALL NODES ##
You can list all nodes using +all+ or +list+

nodes.all #=> [<Chef::Node...>, <Chef::Node...>, ...]

To limit the information returned for each node, pass a code block to the +all+
subcommand:

nodes.all { |node| node.name } #=> [NODE1_NAME, NODE2_NAME, ...]
```

Exercise: Look at the company attribute for the node object

```
chef > node['fqdn']
```

```
=> "centos63.example.com"
```

Exercise: Change the context you are in to that of a cookbook attribute file

```
chef > attributes_mode
```

```
chef:attributes >
```

Exercise: Add a normal attribute

```
chef:attributes > node.normal["interactive"] = "is cool"
```

```
=> "is cool"
```

Exercise: Add a normal attribute

```
chef:attributes > node["interactive"]
```

```
=> "is cool"
```

Exercise: Search for new default attribute

```
$ knife node show node1 -a interactive
```

```
node1:  
  interactive:
```

Exercise: Save node state

```
chef:attributes > node.save
```

```
=> <Chef::Node:0x11671f4 @name="node1">
```

Exercise: Search for new default attribute

```
$ knife node show node1 -a interactive
```

```
node1:  
  interactive: is cool
```

Exercise: Display all node attributes

```
$ knife node show node1 -Fj -l |more
```

```
{
  "name": "node1",
  "chef_environment": "_default",
  "run_list": [
    "role[webserver]"
  ],
  "normal": {
    "tags": [
      ],
      "interactive": "is cool"
    },
    "default": {
      "logrotate": {
        "global": {
          "weekly": true,
          "rotate": 4,
          "create": "",
          "/var/log/wtmp": {
            "missingok": true,
            "monthly": true,
            "create": "0664 root utmp",
            }
          }
        }
      }
    }
```

Exercise: Switch to recipe context

```
chef:attributes > recipe_mode
```

```
chef:recipe >
```

Exercise: Use a package resource

```
chef:recipe > package "tree"
```

```
=> <package[tree] @name: "tree" @noop: nil @before: nil  
@params: {} @provider: nil @allowed_actions:  
[:nothing, :install, :upgrade, :remove, :purge, :reconfig]  
@action: :install @updated: false @updated_by_last_action:  
false @supports: {} @ignore_failure: false @retries: 0  
@retry_delay: 2 @source_line: "(irb#2):12:in  
'irb_binding'" @elapsed_time: 0 @candidate_version: nil  
@options: nil @package_name: "tree"  
@resource_name: :package @response_file: nil @source: nil  
@version: nil @cookbook_name: nil @recipe_name: nil>
```

Exercise: Run Chef

`chef:recipe > run_chef`

```
[2014-05-13T12:20:27-04:00] INFO: Processing package[tree] action install ((irb#1) line 2)
[2014-05-13T12:20:27-04:00] DEBUG: package[tree] checking yum info for tree
[2014-05-13T12:20:27-04:00] DEBUG: package[tree] installed version: (none) candidate version: 1.5.3-2.el6
[2014-05-13T12:20:27-04:00] INFO: package[tree] installing tree-1.5.3-2.el6 from base repository
[2014-05-13T12:20:27-04:00] DEBUG: Executing yum -d0 -e0 -y install tree-1.5.3-2.el6
[2014-05-13T12:20:32-04:00] DEBUG: ---- Begin output of yum -d0 -e0 -y install tree-1.5.3-2.el6 ----
[2014-05-13T12:20:32-04:00] DEBUG: STDOUT:
[2014-05-13T12:20:32-04:00] DEBUG: STDERR:
[2014-05-13T12:20:32-04:00] DEBUG: ---- End output of yum -d0 -e0 -y install tree-1.5.3-2.el6 ----
[2014-05-13T12:20:32-04:00] DEBUG: Ran yum -d0 -e0 -y install tree-1.5.3-2.el6 returned 0
[2014-05-13T12:20:32-04:00] INFO: apache_vhost[lions] sending restart action to service[httpd] (delayed)
[2014-05-13T12:20:32-04:00] INFO: Processing service[httpd] action restart (apache::default line 14)
[2014-05-13T12:20:32-04:00] DEBUG: service[httpd] supports status, running
httpd (pid 49825) is running...
[2014-05-13T12:20:32-04:00] DEBUG: service[httpd] is running
httpd      0:off    1:off    2:on 3:on 4:on 5:on 6:off
Stopping httpd: [ OK ]
Starting httpd: [ OK ]
[2014-05-13T12:20:35-04:00] INFO: service[httpd] restarted
=> true
```

Exercise: Examine the resources it created

```
chef:recipe > resources
```

```
...  "apache_vhost[welcome]",
"apache_vhost[lions]", "apache_vhost[clowns]",
"apache_vhost[bears]", "package[tree]" ]
```

Exercise: Grab a single resource

```
chef:recipe > resources("user[frank]")
```

```
<user[frank] @name: "frank" @noop: nil @before: nil @params: {} @provider: nil @allowed_actions: [:nothing, :create, :remove, :modify, :manage, :lock, :unlock] @action: :create @updated: false @updated_by_last_action: false @supports: {::manage_home=>false, ::non_unique=>false} @ignore_failure: false @retries: 0 @retry_delay: 2 @source_line: "/var/chef/cache/cookbooks/users/recipes/default.rb:11:in `block in from_file'" @elapsed_time: 0.001173478 @resource_name: :user @username: "frank" @comment: "Frank Belson" @uid: 2001 @gid: 0 @home: "/home/frank" @shell: "/bin/bash" @password: nil @system: false @manage_home: false @non_unique: false @cookbook_name: "users" @recipe_name: "default">
 => <user[frank] @name: "frank" @noop: nil @before: nil @params: {} @provider: nil @allowed_actions: [:nothing, :create, :remove, :modify, :manage, :lock, :unlock] @action: :create @updated: false @updated_by_last_action: false @supports: {::manage_home=>false, ::non_unique=>false} @ignore_failure: false @retries: 0 @retry_delay: 2 @source_line: "/var/chef/cache/cookbooks/users/recipes/default.rb:11:in `block in from_file'" @elapsed_time: 0.001173478 @resource_name: :user @username: "frank" @comment: "Frank Belson" @uid: 2001 @gid: 0 @home: "/home/frank" @shell: "/bin/bash" @password: nil @system: false @manage_home: false @non_unique: false @cookbook_name: "users" @recipe_name: "default">
```

Exercise: Show the resource *as you would write it if you were very, very verbose*

```
chef:recipe > puts resources("user[frank]").to_text
```

```
<user[frank] @name: "frank" @noop: nil @before: nil @params: {} @provider: nil @allowed_actions:  
[:nothing, :create, :remove, :modify, :manage, :lock, :unlock] @action: :create @updated: false @updated_by_last_action: false  
@supports: {:manage_home=>false, :non_unique=>false} @ignore_failure: false @retries: 0 @retry_delay: 2 @source_line: "/var/chef/  
cache/cookbooks/users/recipes/default.rb:11:in `block in from_file'" @elapsed_time: 0.001173478 @resource_name: :user @username:  
"frank" @comment: "Frank Belson" @uid: 2001 @gid: 0 @home: "/home/frank" @shell: "/bin/bash" @password: nil @system: false  
@manage_home: false @non_unique: false @cookbook_name: "users" @recipe_name: "default">  
# Declared in /var/chef/cache/cookbooks/users/recipes/default.rb:11:in `block in from_file'  
  
user("frank") do  
  action :create  
  supports {:manage_home=>false, :non_unique=>false}  
  retries 0  
  retry_delay 2  
  username "frank"  
  comment "Frank Belson"  
  uid 2001  
  gid 0  
  home "/home/frank"  
  shell "/bin/bash"  
  cookbook_name "users"  
  recipe_name "default"  
end  
=> nil
```

Exercise: Exit to the main context

```
chef:recipe > exit
```

```
chef >
```

Exercise: reset

```
chef > reset
```

```
Loading.....INFO: Run List is [role[webserver]]  
INFO: Run List expands to [motd, users, ntp,  
apache]  
...INFO: Loading cookbooks [apache, motd, ntp,  
pci, users]  
.done.
```

```
=> nil
```

Exercise: List all the nodes

```
chef > nodes.list
```

```
=> [node[node1]]
```

Exercise: Print all the node names and their environment

```
chef > nodes.all { |n| "#{n.name}:#{n.chef_environment}"}
```

```
=> [ "node1:_default" ]
```

Exercise: From your workstation, use knife to do the same thing

```
$ knife exec -E 'nodes.transform(:name => "node1") { |n| n.chef_environment("production")}'
```

Has...no output >:)

Exercise: Check it from search ...

```
$ knife search node "*:*" -a chef_environment
```

```
1 items found

node1:
  chef_environment:  production
```

Exercise: Check it from chef-shell

```
chef > nodes.all { |n| "#{n.name}:#{n.chef_environment}"}
```

```
=> [ "node1:production" ]
```

Best Practice: Be Careful

- Changes are permanent
- You can do the same tricks on roles, data bags (and items), clients... all kinds of stuff
- You can (and will) blow off your foot if you aren't careful

Exercise: Change our environment back to '`_default`'

```
$ knife node environment set node1 _default
```

```
node1:  
  chef_environment: _default
```

Exercise: knife raw to confirm

```
$ knife raw https://api.opscode.com/organizations/  
yourorg/nodes/node1 | grep chef_environment
```

```
"chef_environment": "_default",
```

Review Questions

- What are chef-shell's three startup modes?
- What are chef-shell's three run-time contexts?
- How do you save attribute changes back to the node object?
- How do you invoke a chef run from within chef-shell?
- How would you list nodes from within chef-shell?
- What knife command can be used to invoke operations that you perform on the node object in the chef-shell?

Writing Ohai Plugins

Getting more data from your systems

v1.1.4

Lesson Objectives

- After completing the lesson, you will be able to
 - Explain the function Ohai and Ohai Plugins
 - Explain what's new in Ohai 7?
 - Write your own Ohai Plugins
 - Force Ohai Plugins to run
 - Disable Ohai Plugins

Ohai!

- Ohai writes automatic attributes - they can't be changed by other components of Chef
- Automatic attributes form the base for every node object at the beginning of every Chef run
- Many plugins are enabled by default, but can be disabled if desired

Running Ohai

```
chef@node1:~$ ohai
```

```
{
  "languages": {
    "ruby": {
      "platform": "x86_64-darwin13.0.0",
      "version": "1.9.3",
      "release_date": "2013-11-22",
    ...
  }
```

Ohai Data Collection

- Platform details: redhat, windows, etc.
- Processor information
- Kernel data
- FQDNs
- Cloud provider data (EC2, Azure, Rackspace, etc)
- Windows device drivers
- and more...

What's new in Ohai 7?

- New DSL
 - **collect_data()** - block of Ruby code called by Ohai when it runs
 - **depends** - comma-separated list of attributes collected by another plugin
 - **provides** - comma-separated list of attributes defined by a plugin
- In Chef 11.12.0 or newer

Check Ohai Version

```
chef@node1:~$ ohai -v
```

```
Ohai: 7.0.2
```

Why Write Ohai Plugins?

- Collect specialized data not “in the box”
- Examples:
 - Information about some daemon you want to expose
 - Hardware inventory information from external sources (BMC, IPMI, Remedy ARS?)
 - Statistics collection on-the-cheap (e.g. shove netstat attributes into automatic data)

Writing Ohai Plugins: Apache Modules

- **Problem:** We want to capture all installed Apache modules as part of our node data
- **Success:** Apache module data is visible in our node attributes

Writing Ohai Plugins: apache.modules

```
chef@node1:~$ apachectl -t -D DUMP_MODULES
```

```
Loaded Modules:  
core_module (static)  
mpm_prefork_module (static)  
http_module (static)  
so_module (static)  
auth_basic_module (shared)  
auth_digest_module (shared)  
authn_file_module (shared)  
authn_alias_module (shared)  
...  
...
```

Installing Ohai Plugins

- Custom Ohai plugins are installed with the Ohai cookbook
 - Configures plugin path for Ohai
 - Delivers plugin to Ohai plugins directory with Chef
 - <http://supermarket.chef.io/cookbooks/ohai>
- <http://docs.chef.io/ohai.html>

Exercise: Download the Ohai Cookbook

```
$ knife cookbook site download ohai 2.0.0
```

Downloading ohai from the cookbooks site at
version 2.0.0 to /Users/YOU/chef-repo/
cookbooks/ohai-2.0.0.tar.gz

Cookbook saved: /Users/YOU/chef-repo/
cookbooks/ohai-2.0.0.tar.gz

Exercise: Download the Ohai Cookbook

```
$ tar -zxvf ohai-2.0.0.tar.gz -C cookbooks/
```

```
x ohai/
x ohai/CHANGELOG.md
x ohai/README.md
x ohai/attributes
x ohai/attributes/default.rb
x ohai/files
x ohai/files/default
x ohai/files/default/plugins
x ohai/files/default/plugins/README
x ohai/metadata.json
x ohai/metadata.rb
x ohai/providers
x ohai/providers/hint.rb
x ohai/recipes
x ohai/recipes/default.rb
x ohai/resources
x ohai/resources/hint.rb
```

Exercise: Upload the Ohai Cookbook

```
$ knife cookbook upload ohai
```

```
Uploading ohai [2.0.0]  
Uploaded 1 cookbook.
```

Exercise: Inspect the Ohai Cookbook



OPEN IN EDITOR: cookbooks/ohai/attributes/default.rb

```
# FHS location would be /var/lib/chef/ohai_plugins or similar.
case node["platform_family"]
when "windows"
  default["ohai"]["plugin_path"] = "C:/chef/ohai_plugins"
else
  default["ohai"]["plugin_path"] = "/etc/chef/ohai_plugins"
end

# The list of plugins and their respective file locations
default["ohai"]["plugins"]["ohai"] = "plugins"
```

Exercise: Update apache cookbook

- We could create our modules.rb file and add it to the ohai cookbook, but...
- We now have to maintain a forked version of the community cookbook
- Let's use our apache cookbook instead!

Exercise: Update the metadata.rb



OPEN IN EDITOR: cookbooks/apache/metadata.rb

```
name                  'apache'
maintainer           'Your Name'
maintainer_email     'your\_email@example.com'
license               'All rights reserved'
description          'Installs/Configures apache'
long_description    IO.read(File.join(File.dirname(__FILE__),
'README.md'))
version              '0.4.0'
depends              'ohai'
```

SAVE FILE!

181

Basic Ohai DSL

- **Ohai.plugin(:Name)** - used to identify the plugin

```
Ohai.plugin(:Name) do
  provides "attribute_a", "attribute_b"
  depends "attribute_x", "attribute_y"

  collect_data(:default) do
    # some Ruby code
    attribute_a Mash.new
  end

  collect_data(:platform...) do
    # platform-specific Ruby code
    attribute_b Mash.new
  end
end
```

Basic Ohai DSL

- **Ohai.plugin(:Name)** - used to identify the plugin
- **provides** - comma-separated list of attributes defined by the plugin

```
Ohai.plugin(:Name) do
  provides "attribute_a", "attribute_b"
  depends "attribute_x", "attribute_y"

  collect_data(:default) do
    # some Ruby code
    attribute_a Mash.new
  end

  collect_data(:platform...) do
    # platform-specific Ruby code
    attribute_b Mash.new
  end
end
```

Basic Ohai DSL

- **Ohai.plugin(:Name)** - used to identify the plugin
- **provides** - comma-separated list of attributes defined by the plugin
- **depends** - comma-separated list of attributes collected by another plugin

```
Ohai.plugin(:Name) do
  provides "attribute_a", "attribute_b"
  depends "attribute_x", "attribute_y"

  collect_data(:default) do
    # some Ruby code
    attribute_a Mash.new
  end

  collect_data(:platform...) do
    # platform-specific Ruby code
    attribute_b Mash.new
  end
end
```

Basic Ohai DSL

- **Ohai.plugin(:Name)** - used to identify the plugin
- **provides** - comma-separated list of attributes defined by the plugin
- **depends** - comma-separated list of attributes collected by another plugin
- **collect_data(:default)** - the default code run if the platform is not defined

```
Ohai.plugin(:Name) do
  provides "attribute_a", "attribute_b"
  depends "attribute_x", "attribute_y"

  collect_data(:default) do
    # some Ruby code
    attribute_a Mash.new
  end

  collect_data(:platform...) do
    # platform-specific Ruby code
    attribute_b Mash.new
  end
end
```

Basic Ohai DSL

- **Ohai.plugin(:Name)** - used to identify the plugin
- **provides** - comma-separated list of attributes defined by the plugin
- **depends** - comma-separated list of attributes collected by another plugin
- **collect_data(:default)** - the default code run if the platform is not defined
- **collect_data(:platform)** - the code run for a specific platform

```
Ohai.plugin(:Name) do
  provides "attribute_a", "attribute_b"
  depends "attribute_x", "attribute_y"

  collect_data(:default) do
    # some Ruby code
    attribute_a Mash.new
  end

  collect_data(:platform...) do
    # platform-specific Ruby code
    attribute_b Mash.new
  end
end
```

Exercise: Create modules.rb



OPEN IN EDITOR: apache/files/default/plugins/modules.rb

```
Ohai.plugin(:Apache) do
  require 'mixlib/shellout'
  provides "apache/modules"

  collect_data(:default) do
    modules = Mixlib::ShellOut.new("apachectl -t -D DUMP_MODULES")
    modules.run_command
    apache Mash.new
    apache[:modules] = modules.stdout
  end
end
```

SAVE FILE !

187

Mixlib::Shellout

- Cross-platform library for safely executing shell commands from within Chef, Ohai, etc.
- Deals with passing arguments, capturing STDERR and STDOUT, timing-out hung processes
- Deals with cross-platform (Linux, UNIX, Windows) subprocess quirks
- **Never use backticks/Process.spawn/system!**
- **Use Mixlib::Shellout!**

Exercise: Write the ohai_plugin recipe



OPEN IN EDITOR: cookbooks/apache/recipes/ohai_plugin.rb

```
ohai 'reload_apache' do
  plugin 'apache'
  action :nothing
end

cookbook_file "#{node['ohai']['plugin_path']}/modules.rb" do
  source 'plugins/modules.rb'
  owner 'root'
  group 'root'
  mode '0644'
  notifies :reload, 'ohai[reload_apache]', :immediately
end

include_recipe 'ohai::default'
```

Exercise: Upload the apache cookbook

```
$ knife cookbook upload apache
```

```
Uploading apache [ 0 . 4 . 0 ]  
Uploaded 1 cookbook.
```

Exercise: Add apache::ohai_plugin to the run list

```
$ knife node run_list set node1  
'recipe[apache::ohai_plugin]', 'role[webserver]'
```

- For Powershell you need to do this

```
$ knife node run_list set node1  
`'recipe[apache::ohai_plugin],role[webserver]`'
```

```
node1:  
  run_list:  
    recipe[apache::ohai_plugin]  
    role[webserver]
```

Exercise: Run chef-client

```
chef@node1:~$ sudo chef-client
```

```
Starting Chef Client, version 11.10.4.ohai7.0
resolving cookbooks for run list: ["apache::ohai_plugin"]
Synchronizing Cookbooks:
  - apache
  - ohai
Compiling Cookbooks...
Recipe: ohai::default
* remote_directory[/etc/chef/ohai_plugins] action create
  - create new directory /etc/chef/ohai_plugins
  - change mode from '' to '0755'
  - restore selinux security context
Recipe: <Dynamically Defined Resource>
```

Exercise: Run chef-client

```
* cookbook_file[/etc/chef/ohai_plugins/README] action create
- create new file /etc/chef/ohai_plugins/README
- update content in file /etc/chef/ohai_plugins/README from none to 775fa7
  --- /etc/chef/ohai_plugins/README    2014-02-27 09:31:04.558032870 -0500
  +++ /tmp/.README20140227-33098-tpmpwc  2014-02-27 09:31:04.621065128
-0500
  @@ -1,2 @@
  +This directory contains custom plugins for Ohai.
- change mode from '' to '0644'
- restore selinux security context

Recipe: ohai::default
* ohai[custom_plugins] action reload/opt/chef/embedded/lib/ruby/gems/1.9.1/
gems/ohai-7.0.0.rc.0/lib/ohai/plugins/linux/network.rb:49: warning: already
initialized constant IPROUTE_INT_REGEX

- re-run ohai and merge results into node attributes
```

Exercise: Run chef-client

```
Converging 3 resources
Recipe: apache::ohai_plugin
 * cookbook_file[/etc/chef/ohai_plugins/modules.rb] action create
- create new file /etc/chef/ohai_plugins/modules.rb
- update content in file /etc/chef/ohai_plugins/modules.rb from none to 2d1a51
  --- /etc/chef/ohai_plugins/modules.rb 2014-02-27 09:31:09.667647850 -0500
  +++ /tmp/.modules.rb 2014-02-27 09:31:09.718673962 -0500
  @@ -1 +1,9 @@
+Ohai.plugin(:Apache) do
+  provides "apache/modules"
+
+  collect_data(:default) do
+    modules = shell_out("apachectl -t -D DUMP_MODULES")
+    apache Mash.new
+    apache[:modules] = modules
+  end
+end
- change mode from '' to '0755'
- change owner from '' to 'root'
- change group from '' to 'root'
- restore selinux security context
```

Exercise: Run chef-client

```
Recipe: ohai::default
  * ohai[custom_plugins] action reload/opt/chef/embedded/lib/ruby/gems/
    1.9.1/gems/ohai-7.0.0.rc.0/lib/ohai/plugins/linux/network.rb:49: warning:
    already initialized constant IPROUTE_INT_REGEX

    - re-run ohai and merge results into node attributes

  * remote_directory[/etc/chef/ohai_plugins] action nothing (skipped due to
    action :nothing)
  * ohai[custom_plugins] action nothing (skipped due to action :nothing)

Running handlers:
Running handlers complete

Chef Client finished, 7/53 resources updated in 16.185836857 seconds
```

Exercise: Show apache attribute

```
$ knife node show node1 -a apache
```

```
node1:  
apache:  
  indexfile: index1.html  
  modules:  Loaded Modules:  
    core_module (static)  
    mpm_prefork_module (static)  
    http_module (static)  
    so_module (static)  
    auth_basic_module (shared)  
  ...
```

Exercise: View plugin

```
chef@node1:~$ cat /etc/chef/ohai_plugins/modules.rb
```

```
Ohai.plugin(:Apache) do
  require 'mixlib/shellout'
  provides "apache/modules"

  collect_data(:default) do
    modules = Mixlib::ShellOut.new("apachectl -t -D DUMP_MODULES")
    modules.run_command
    apache Mash.new
    apache[:modules] = modules.stdout
  end
end
```

Can We Make It Better?

- We have Apache module data!
- It's saved as one big string... that's a bummer
- Let's sort this into a group of static and shared modules

Exercise: Refactor modules.rb



OPEN IN EDITOR: apache/files/default/plugins/modules.rb

```
Ohai.plugin(:Apache) do
  require 'mixlib/shellout'
  provides "apache/modules"

  collect_data(:default) do
    apache Mash.new
    apache[:modules] = { :static => [], :shared => [] }
```

- Let's start by setting up our node['apache'] Mash

Exercise: Refactor modules.rb



OPEN IN EDITOR: apache/files/default/plugins/modules.rb

```
modules = shell_out("apachectl -t -D DUMP_MODULES")
modules.stdout.each_line do |line|
  fullkey, value = line.split(/\(/, 2).map { |i| i.strip}
  apache_mod = fullkey.gsub(/_module/, "")
  dso_type = value.to_s.chomp("\\")

  if dso_type.match(/shared/)
    apache[:modules][:shared] << apache_mod
  elsif dso_type.match(/static/)
    apache[:modules][:static] << apache_mod
  end
end
end
end
```

SAVE FILE!

200

Exercise: Upload the apache cookbook

```
$ knife cookbook upload apache
```

```
Uploading apache [ 0 . 4 . 0 ]  
Uploaded 1 cookbook.
```

Exercise: Run chef-client

```
chef@node1:~$ sudo chef-client
```

```
Starting Chef Client, version 11.12.8
resolving cookbooks for run list: ["apache::ohai_plugin"]
Synchronizing Cookbooks:
  - apache
  - ohai
Compiling Cookbooks...
Recipe: apache::ohai_plugin
 * ohai[reload_apache] action nothing (skipped due to action :nothing)
 * cookbook_file[/etc/chef/ohai_plugins/modules.rb] action create
   - update content in file /etc/chef/ohai_plugins/modules.rb from e6cf9a to
809f46
   --- /etc/chef/ohai_plugins/modules.rb 2014-04-02 23:22:30.727649249
-0400
```

Exercise: Show apache.modules attributes

```
$ knife node show node1 -a apache.modules
```

```
node1:  
apache.modules:  
  shared:  
    auth_basic  
    auth_digest  
    authn_file  
    authn_alias  
    authn_anon  
  ...  
  static:  
    core  
    mpm_prefork  
    http  
    so
```

Exercise: Start the chef-shell REPL

```
chef@node1$ /opt/chef/bin/chef-shell
```

```
loading configuration: none (standalone session)
Session type: standalone
Loading.....done.
```

```
This is the chef-shell.
Chef Version: 11.12.8
http://www.opscode.com/chef
http://docs.opscode.com/
```

```
run `help` for help, `exit` or ^D to quit.
```

```
Ohai2u chef@centos63.example.com!
chef >
```

Exercise: Test Ohai Plugins with chef-shell

```
chef > Ohai::Config['plugin_path'] << "/etc/chef/
ohai_plugins"
=> ["/etc/chef/ohai_plugins", "/opt/chef/embedded/lib/
ruby/gems/1.9.1/gems/ohai-7.0.4/lib/ohai/plugins", "/etc/
chef/ohai_plugins"]
```

Exercise: Test Ohai Plugins with chef-shell

```
chef > o = Ohai::System.new  
  
=> #<Ohai::System:0x0000000186acd0 @plugin_path="",  
@data={}, @provides_map=#<Ohai::ProvidesMap:  
0x0000000186ac30 @map={}, @v6_dependency_solver={},  
@loader=#<Ohai::Loader:0x0000000186ab90  
@controller=#<Ohai::System:0x0000000186acd0 ...>,  
@v6_plugin_classes[], @v7_plugin_classes[],  
@runner=#<Ohai::Runner:0x0000000186ab18  
@provides_map=#<Ohai::ProvidesMap:0x0000000186ac30  
@map={}, @safe_run=true>>
```

Exercise: Test Ohai Plugins with chef-shell

```
chef > o.all_plugins  
=> [ .....]
```

Exercise: Test Ohai Plugins with chef-shell

```
chef > pp o.refresh_plugins(attribute_filter=
'apache/modules')

=> [#<Ohai::NamedPlugin::Apache:0x00000001f602d8
@data={"apache"=>{"modules"=>{"static"=>["core",
"mpm_prefork", "http", "so"],
"shared"=>["auth_basic", ...}}
```

Plugin Debugging Notes

- Ohai will **ignore plugin failures/exceptions**
- This is to avoid breaking Chef Client during such an early phase
- If you get no data from your plugins, this is probably why
- **Write plugins defensively. Check for all error conditions!**

Ohai Hints

- Ohai plugins can be forced to run, even if detected system state says that they shouldn't!
- For example, it's impossible to reliably detect whether a server is in EC2, so EC2 plugin may not run
- Solution: Force the EC2 plugin to run by placing an empty JSON document ("{}") in
`/etc/chef/ohai/hints/ec2.json`

Disabling Ohai Plugins

- Some plugins can execute slowly on different platforms
- You can disable plugins whose data you don't use to save speed things up!
- Only disables plugins for Ohai when run with chef-client

The Problem & Success Criteria

- **The Problem:** Our systems are LDAP/Active Directory joined and thus the Ohai data contains thousands of user accounts, slowing down the Chef run.
- **Success Criteria:** We will disable the :Passwd plugin to avoid collecting this information

Exercise: Disable :Passwd plugin

```
$ knife node show node1 -a etc.passwd
```

```
node1:  
etc.passwd:  
abrt:  
  dir: /etc/abrt  
  gecos:  
  gid: 173  
  shell: /sbin/nologin  
  uid: 173  
adm:  
  dir: /var/adm  
  gecos: adm  
  gid: 4  
  shell: /sbin/nologin  
  uid: 3  
...
```

Exercise: Update the base role



OPEN IN EDITOR: roles/base.rb

```
name "base"
description "Base Server Role"
run_list "recipe[chef-client::delete_validation]", "recipe[chef-client::config]", "recipe[chef-client]", "recipe[ntp]", "recipe[motd]", "recipe[users]"
default_attributes(
  "chef_client" => {
    "config" => {
      "ssl_verify_mode" => ":verify_peer"
      "log_level" => ":info"
    }
  },
  "ohai" => {
    "disabled_plugins" => [":Passwd"]
  }
)
```

-214-

Exercise: Upload the base role

```
$ knife role from file base.rb
```

Updated Role base!

Exercise: Disable :Passwd plugin

```
chef@node1:~$ sudo chef-client && sudo chef-client
```

```
Starting Chef Client, version 11.12.4
resolving cookbooks for run list: ["apache::ohai_plugin",
"chef-client::delete_validation", "chef-client", "ntp",
"motd", "users", "apache"]
...
Running handlers:
Running handlers complete

Chef Client finished, 1/4 resources updated in 9.690084663
seconds
```

Exercise: Disable :Passwd plugin

```
$ knife node show node1 -a etc.passwd
```

```
node1:  
  etc.passwd:
```

Exercise: Disable passwd plugin

```
chef@node1:~$ sudo cat /etc/chef/client.rb
```

```
chef_server_url "https://api.opscode.com/organizations/intermediate050614"
validation_client_name "intermediate050614-validator"
verify_api_cert true
ssl_verify_mode :verify_peer
node_name "node1"
```

```
Ohai::Config[:plugin_path] << "/etc/chef/ohai_plugins"
Ohai::Config[:disabled_plugins] = [:Passwd]
```

```
Dir.glob(File.join("/etc/chef", "client.d", "*.rb")).each do |conf|
  Chef::Config.from_file(conf)
end
```

Other Tips

- Extreme situations: whitelist only the attributes you want to send: <http://ckbk.it/whitelist-node-attrs>
 - Minimize the amount of JSON you are sending across the wire on every run
 - Useful if you have many (>10,000) nodes
- Docs: <http://docs.chef.io/ohai.html>

Special Upgrading Note

- The syntax for disabling Ohai plugins has changed slightly from 6 to 7
- **Ohai 6**
 - `Ohai::Config[:disabled_plugins] = ["passwd"]`
- **Ohai 7**
 - `Ohai::Config[:disabled_plugins] = [:Passwd]`

Review Questions

- What command can you run at the command line to invoke Ohai?
- How are Ohai plugins installed?
- What is wrong with this first line of an Ohai definition: "Ohai.plugin(:name) do"?

Chef Client Run Internals

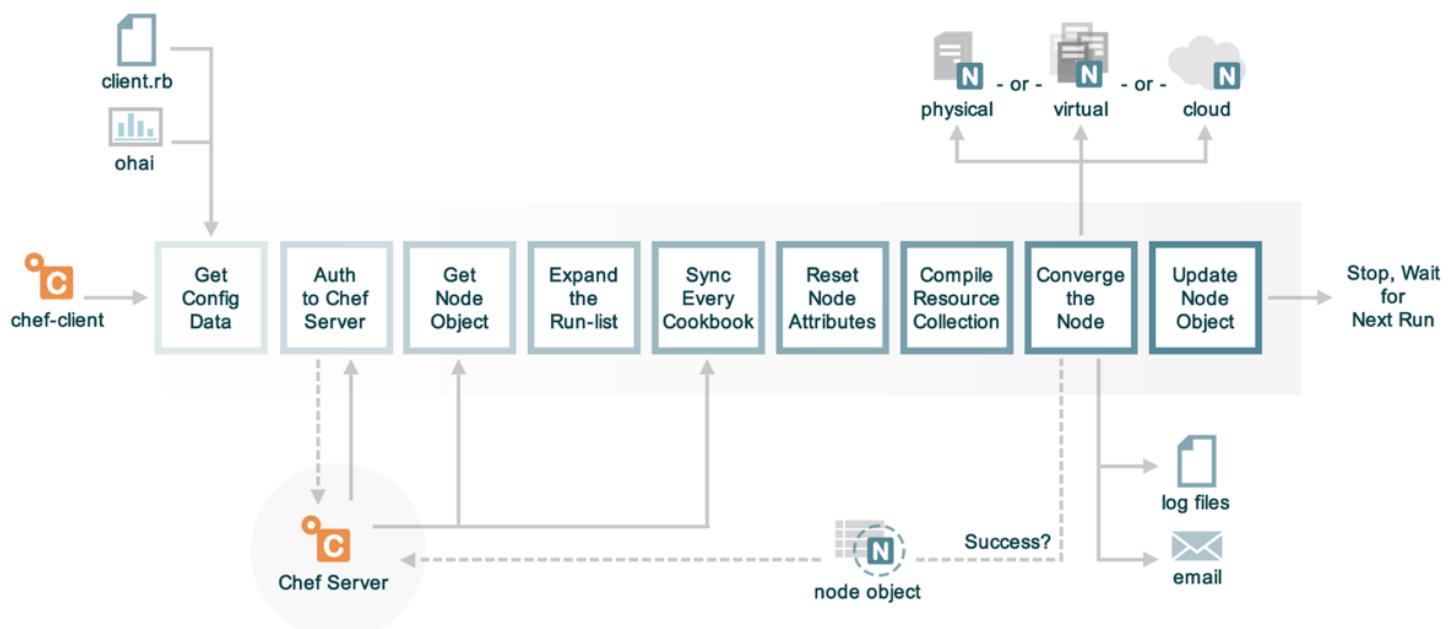
What really happens when Chef runs?

v1.1.4

Lesson Objectives

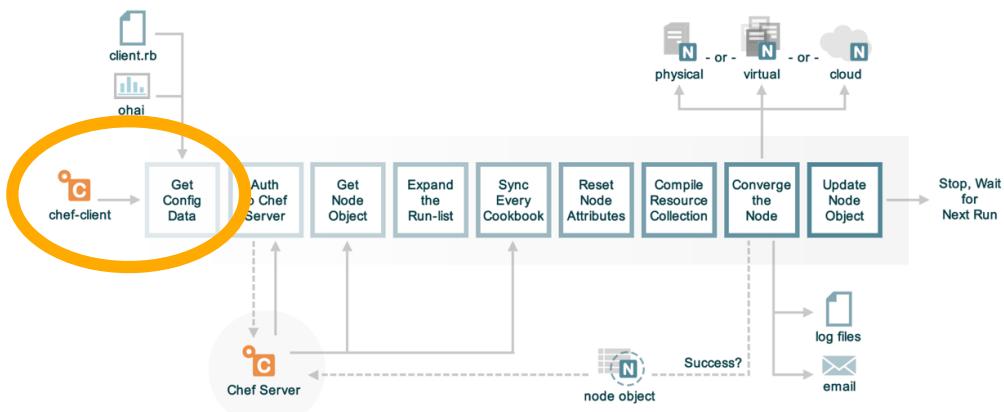
- After completing the lesson, you will be able to:
 - Understand what happens under the hood when Chef Client runs
 - Describe the purpose of the `run_context` and `run_status` objects
 - Use `node.run_state` to pass information around the recipe at execution time
 - Describe the event subsystem and how it can be used

Remember This From Fundamentals?



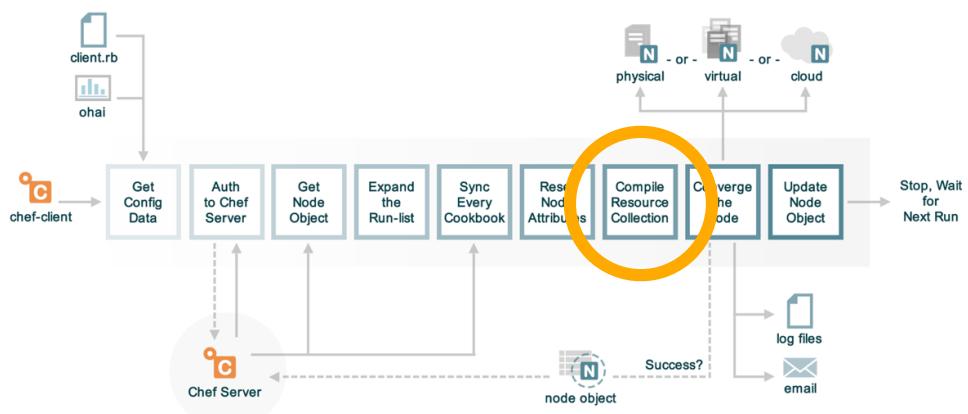
Customizations

- Ohai Plugins
- Start Handlers
- Event Dispatcher



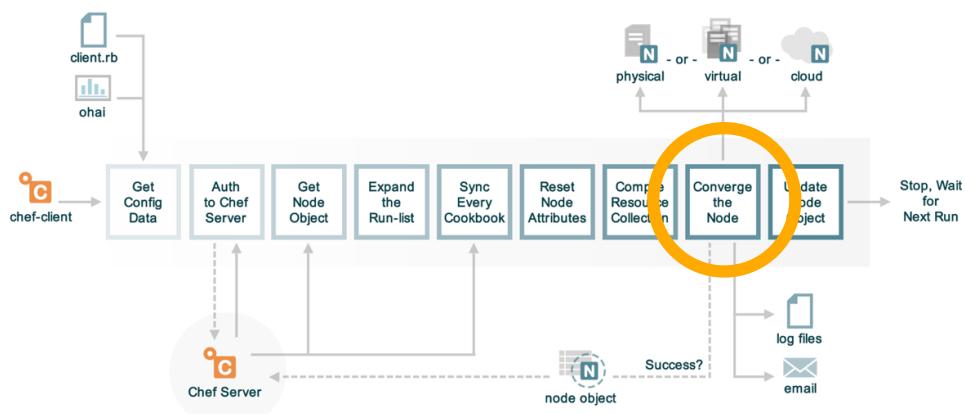
Customizations

- LWRPs
- HWRPs
- Libraries



Customizations

- Exception Handlers
- Report Handlers



What is the run_context?

- Master object for the Chef run for this node
- Important parts of the hierarchy:
 - `run_context`
 - `node`
 - `chef_environment`
 - `run_status`
 - `run_state`
 - `cookbook_collection`
 - `resource_collection`
 - `events`

What is the run_status?

- Hangs off node
- Stores current status of run
- Resources that are converged, resources that changed during this run
- Run status methods (`success?` / `failed?`)
- We'll use this object later when we write event handlers

What is the run_state?

- A Hash that hangs off node
- A place for you to stash transient data during the run
- Transmits data between resources



The Problem and Success Criteria

- **Problem:** We can't decide what scripting language we want to use to write our web application.
- **Success Criteria:** Chef has randomly chosen one for us each time it runs.



or



Practical Extraction and Report Language

Exercise: Edit Apache recipe



OPEN IN EDITOR: apache/recipes/default.rb

```
package "httpd" do
  action :install
end

ruby_block "randomly_choose_language" do
  block do
    if Random.rand > 0.5
      node.run_state['scripting_language'] = 'php'
    else
      node.run_state['scripting_language'] = 'perl'
    end
  end
end

package "scripting_language" do
  package_name lazy { node.run_state['scripting_language'] }
  action :install
end
```

SAVE FILE !

232

ruby_block resource

```
ruby_block "randomly_choose_language" do
  block do
    if Random.rand > 0.5
      node.run_state['scripting_language'] = 'php'
    else
      node.run_state['scripting_language'] = 'perl'
    end
  end
end
```

- `ruby_block` declares a block of Ruby to run at **execute** time
- It has direct access to the `node` structure and other aspects of the run
- Store something on `node.run_state` for later use

node.run_state

```
node.run_state[ 'scripting_language' ] = 'php'
```

- `node.run_state` is an empty Hash
- Discarded at the end of the run

Lazy parameter evaluation

```
package "scripting_language" do
  package_name lazy { node.run_state['scripting_language'] }
  action :install
end
```

- Normally resource parameters must be specified completely at **compile** time
- Sometimes their value is not known until **execute** time
- Use the `lazy{ }` block to have parameters evaluated then

Upload the new apache cookbook

```
$ knife cookbook upload apache
```

```
Uploading apache [0.4.0]
Uploaded 1 cookbook.
```

Run Chef Client

```
chef@node1$ sudo chef-client
```

```
* ruby_block[randomly_choose_language] action run
  - execute the ruby block randomly_choose_language

* package[scripting_language] action install
  - install version 5.3.3-27.el6_5 of package php
```

Event Dispatcher Subsystem

- Used by logging, Enterprise Chef reporting
- Simple queue-free publish/subscribe model
- All listeners get all events
- Log formatters (:min, :doc, etc.) are examples of Event Dispatcher subclasses

Implementing Event Dispatchers

- Write a subclass of `Chef::EventDispatch::Base`
 - Base methods are all set to take no action
- Implement methods in subclass to be called
 - Examples: `run_start`, `run_completed`
- Create a start handler to load your event dispatcher
- Register your handler with Chef (next chapter)

Review Questions

- What is the name of the root object of the Chef run?
- Where would you stash transient data during the run?
- What object stores the current status of the run?
- What can be used to populate resource attributes at execute time?

Implementing Chef Handlers

Communicating the status of a Chef Client run

v1.1.4

Lesson Objectives

- After completing this lesson, you will be able to
 - Describe Chef handlers
 - Write custom report handlers
 - Distribute handlers to nodes
 - Configure Chef to use custom handlers

Handlers

- Handlers run at the start or end of a chef-client run
- They are *Ruby programs* that can read the status information of your chef-client run
- Three types:
 - Report
 - Exception
 - Start

Report Handlers

- Used when a chef-client run succeeds
- Runs when the `run_status.success?` is `true`
- Community Examples:
 - http://ampledata.org/splunk_storm_chef_handler.html
 - <http://onddo.github.io/chef-handler-zookeeper/>
 - https://github.com/realityforge/chef-graphite_handler

Exception Handlers

- Used when a chef-client run fails
- Runs when the `run_status.failed?` is `true`
- Community Examples:
 - https://github.com/morgoth/airbrake_handler
 - https://github.com/jblaine/syslog_handler
 - <http://onddo.github.io/chef-handler-sns/>

Start Handlers

- Used to run events at the beginning of the chef-client run
- May **NOT** be loaded using the `chef_handler` resource
- Install with `chef_gem` or add to the `start_handlers` setting in the client.rb
- <https://github.com/opscode/chef-reporting>

Writing Custom Handlers

- Use any ruby gem or library
- Distribute to nodes with the **chef_handler** or **chef-client** cookbooks
- Optionally configure through **client.rb** settings

Exercise: Download the chef_handler cookbook

```
$ knife cookbook site download chef_handler 1.1.6
```

```
Downloading chef_handler from the cookbooks site  
at version 1.1.6 to /Users/YOU/chef-repo/  
chef_handler-1.1.6.tar.gz
```

```
Cookbook saved: /Users/YOU/chef-repo/  
chef_handler-1.1.6.tar.gz
```

Exercise: Download the chef-client cookbook

```
$ tar -zxvf chef_handler-1.1.6.tar.gz -C cookbooks/
```

```
x chef_handler/
x chef_handler/CHANGELOG.md
x chef_handler/README.md
x chef_handler/attributes
x chef_handler/attributes/default.rb
x chef_handler/files
x chef_handler/files/default
x chef_handler/files/default/handlers
x chef_handler/files/default/handlers/README
x chef_handler/libraries
x chef_handler/libraries/matchers.rb
```

Exercise: Upload the chef_handler cookbook

```
$ knife cookbook upload chef_handler
```

```
Uploading chef_handler [1.1.6]
Uploaded 1 cookbook.
```

The Problem and Success Criteria

- **Problem:** When a chef-client run ends we want to be notified via email of only the resources that have changed on the node.
- **Success Criteria:** We receive an email containing the resource changed on the node.

Let's Write a Handler

- We're going to write a new handler that will **email** us when chef-client runs
- It will send us all of the resources that changed during the chef-client run
- <http://docs.chef.io/handlers.html>

Dear Ruby, How Do You Email?

- Handlers are *Ruby programs* that can read the status information of your chef-client run.
- We could write a library to send an email or we could look for one:
 - Ruby's Standard Library
 - Ruby's Extended Library
 - Ruby gems

Dear Ruby, How Do You Email?

- <http://www.rubydoc.info/stdlib/core/1.9.3>
- <http://www.rubydoc.info/stdlib>
- <https://www.rubygems.org/>
- https://www.ruby-toolbox.com/categories/e_mail

Library Cookbook Pattern

- We could modify the **chef_handler** cookbook to add a recipe for the email handler
- However, we'd basically be “**forking**” an upstream cookbook
- Instead, let's create our own cookbook and use **chef_handler** as a “**library**”
- Code reusability!

Exercise: Create a cookbook named ‘email_handler’

```
$ knife cookbook create email_handler
```

```
** Creating cookbook email_handler
** Creating README for cookbook: email_handler
** Creating CHANGELOG for cookbook: email_handler
** Creating metadata for cookbook: email_handler
```

Exercise: Edit the default recipe



OPEN IN EDITOR: cookbooks/email_handler/recipes/default.rb

```
chef_gem "pony" do
  action :install
end

include_recipe "chef_handler"
```

SAVE FILE !

257

The `chef_handler` Resource

- `chef_handler` is a resource packaged with the `chef_handler` cookbook
- It has two actions, `:enable` and `:disable`
- It has three arguments
 - `source` - the file containing the handler code
 - `arguments` - any pieces of information needed to initialize the handler
 - `supports` - `:report`, `:exception`
- Defaults:
 - `:enable`
 - `:report => true`, `:exception => true`

Exercise: Setup the handler



OPEN IN EDITOR: cookbooks/email_handler/recipes/default.rb

```
chef_gem "pony" do
  action :install
end

include_recipe "chef_handler"

cookbook_file "#{node['chef_handler']['handler_path']}/email_handler.rb" do
  source "handlers/email_handler.rb"
  owner "root"
  group "root"
  mode "0644"
end
```

SAVE FILE !

259

Exercise: Setup the handler



OPEN IN EDITOR: cookbooks/email_handler/recipes/default.rb

```
cookbook_file "#{node['chef_handler']['handler_path']}/email_handler.rb" do
  source "handlers/email_handler.rb"
  owner "root"
  group "root"
  mode "0644"
end

chef_handler "MyCompany::EmailMe" do
  source "#{node['chef_handler']['handler_path']}/email_handler.rb"
  arguments [node['email_handler']['from_address'],
             node['email_handler']['to_address']]
  action :enable
end
```

SAVE FILE !

260

Exercise: Set the attributes



OPEN IN EDITOR: cookbooks/email_handler/attributes/default.rb

```
default['email_handler']['from_address'] = "chef@localhost"
default['email_handler']['to_address'] = "chef@localhost"
```

SAVE FILE!

Exercise: Write the handler



OPEN IN EDITOR: .../email_handler/files/default/handlers/email_handler.rb

```
require 'rubygems'  
require 'pony'  
  
module MyCompany  
  class EmailMe < Chef::Handler
```

SAVE FILE!

- All custom exception and report handlers are defined using Ruby and must be a subclass of the **Chef::Handler** class.
- The module and class match what we defined as the name of the `chef_handler` in the recipe

The initialize Method



OPEN IN EDITOR: .../email_handler/files/default/handlers/email_handler.rb

```
class EmailMe < Chef::Handler

  def initialize(from_address, to_address)
    @from_address = from_address
    @to_address = to_address
  end
```

SAVE FILE!

- Initialize the handler with the arguments we passed in the definition
- You can create the method with any args you need to meet your requirements

The report Method



OPEN IN EDITOR: .../email_handler/files/default/handlers/email_handler.rb

```
@to_address = to_address
end

def report
  status = "Failed"
  if success?
    status = "Successful"
  end
  subject = "#{status} Chef run report from #{node.name}"
  report_string = ""      SAVE FILE!
```

- The **report** interface is used to define how a handler will behave and is a required part of any custom handler

The updated_resources Hash



OPEN IN EDITOR: [../email_handler/files/default/handlers/email_handler.rb](#)

```
# report on changed resources
if ! run_status.updated_resources.empty?
  # get some info about all the changed resources!
  run_status.updated_resources.each do |r|
    report_string += "The resource #{r.name} was changed in cookbook
#{r.cookbook_name} at #{r.source_line}\n"
  end
else
  report_string += "No resources changed by chef-client\n"
end
```

SAVE FILE!

- **updated_resources** records information about all the resources changed during a chef-client run
- read through this hash with **.each**, pull interesting information out about each resource

Exercise: Finish email_handler.rb



OPEN IN EDITOR: ../email_handler/files/default/handlers/email_handler.rb

```
report_string += "No resources changed by chef-client\n"
end

Pony.mail(:to => @to_address,
           :from => @from_address,
           :subject => subject,
           :body => report_string)
end
end
end
```

SAVE FILE!

- Use **Pony.mail** to send a message containing info on changed resources

Other Dependencies

- Make sure all necessary functions are available!
 - Some sort of mail transfer agent (MTA)
 - Some sort of email reader (MUA)
- These are distinct functional pieces
 - May have uses other than the handler
 - Get to be their own cookbooks!

Exercise: Download the postfix cookbook

```
$ knife cookbook site download postfix 3.1.8
```

Downloading postfix from the cookbooks site
at version 3.1.8 to /Users/YOU/chef-repo/
postfix-3.1.8.tar.gz

Cookbook saved: /Users/YOU/chef-repo/
postfix-3.1.8.tar.gz

Exercise: Download the postfix cookbook

```
$ tar -zxvf postfix-3.1.8.tar.gz -C cookbooks/
```

```
x postfix/
x postfix/CHANGELOG.md
x postfix/README.md
x postfix/attributes
x postfix/attributes/default.rb
x postfix/files
x postfix/files/default
x postfix/files/default/tests
x postfix/files/default/tests/minitest
x postfix/files/default/tests/minitest/support
```

Exercise: Upload the postfix cookbook

```
$ knife cookbook upload postfix
```

```
Uploading postfix  
upload complete
```

```
[ 3.1.8 ]
```

MUA: mailutils

- There isn't a community cookbook for `mailx`, a command line mail reader
- On Debian-family systems, `mailx` is included with the `mailutils` package
- On Red Hat-family systems, `mailx` is included by the `mailx` package

Exercise: Create the mailx cookbook

```
$ knife cookbook create mailx
```

```
** Creating cookbook mailx
** Creating README for cookbook: mailx
** Creating CHANGELOG for cookbook: mailx
** Creating metadata for cookbook: mailx
```

Exercise: Install the package



OPEN IN EDITOR: cookbooks/mailx/recipes/default.rb

```
package "mailx" do
  action :install
end
```

SAVE FILE !

Exercise: Add the mail dependencies



OPEN IN EDITOR: cookbooks/email_handler/recipes/default.rb

```
chef_gem "pony" do
  action :install
end

include_recipe "chef handler"
include_recipe "postfix"
include_recipe "mailx"

cookbook_file "#{node['chef_handler']['handler_path']}/
email_handler.rb" do
  source "handlers/email_handler.rb"
  owner "root"
```

SAVE FILE !

274

Exercise: Update the metadata.rb



OPEN IN EDITOR: cookbooks/email_handler/metadata.rb

```
name          "email_handler"
maintainer    "You"
maintainer_email "you@somewhere.com"
license        "Apache 2.0"
description    "Email me on every Chef run"
long_description IO.read(File.join(File.dirname(__FILE__),  
'README.md'))
version        "0.1.0"

depends "chef_handler"
depends "postfix"
depends "mailx"
```

SAVE FILE !

Exercise: Upload the email_handler cookbook

```
$ knife cookbook upload email_handler  
postfix mailx
```

```
Uploading email_handler [ 0.1.0 ]  
Uploading postfix [ 3.1.8 ]  
Uploading mailx [ 0.1.0 ]  
Uploaded 3 cookbooks.
```

Exercise: Add email_handler to base role



OPEN IN EDITOR: roles/base.rb

```
name "base"
description "Base Server Role"
run_list "recipe[email_handler]", "recipe[chef-client::delete_validation]",
"recipe[chef-client::config]", "recipe[chef-client]", "recipe[ntp]",
"recipe[motd]", "recipe[users]"
```

SAVE FILE !

277

Best Practice: Add handlers at beginning of run

- Set up handlers (especially exception handlers) at the beginning of the run_list
- That way, if the run fails, the handler will still execute

Exercise: Upload the base role

```
$ knife role from file base.rb
```

Updated Role base!

Exercise: Run chef-client

```
chef@node1$ sudo chef-client
```

```
...
* chef_handler[MyCompany::EmailMe] action enable
  - load /var/chef/handlers/email_handler.rb
    - enable chef_handler[MyCompany::EmailMe] as a
report handler
      - enable chef_handler[MyCompany::EmailMe] as a
exception handler
...
...
```

Read the Message Using "mail"

```
chef@node1$ mail
```

```
Heirloom Mail version 12.4 7/29/08. Type ? for help.
"/var/spool/mail/chef": 1 message 1 new
>N 1 pony@unknown      Wed May 14 09:14 30/2412 "Successful Chef run report from node1"
& r
To: chef@localhost pony@unknown
Subject: Re: Successful Chef run report from node1

pony@unknown wrote:

> The resource pony was changed in cookbook email_handler at /var/chef/cache/cookbooks/
email_handler/recipes/default.rb:9:in `from_file'
> The resource /var/chef/handlers was changed in cookbook chef_handler at /var/chef/cache/
cookbooks/chef_handler/recipes/default.rb:23:in `from_file'
> The resource /etc/postfix/main.cf was changed in cookbook postfix at /var/chef/cache/
cookbooks/postfix/recipes/default.rb:95:in `block in from_file'
> The resource /etc/postfix/master.cf was changed in cookbook postfix at /var/chef/cache/
cookbooks/postfix/recipes/default.rb:95:in `block in from_file'
```

Review Questions

- What cookbook is used to configure Chef Handlers?
- What are the three types of Chef Handler?
- When are Start Handlers run?
- When are Report Handlers run?
- When are Exception Handlers run?

Cookbook Style & Correctness

The real benefits of Infrastructure As Code

v1.1.4

Lesson Objectives

- After completing this lesson, you will be able to:
 - Explain the benefits of correctness checking your cookbooks
 - Explain how and why you should test your cookbooks
 - Use Foodcritic to check the validity of your cookbooks
 - Verify your code adheres to the community Ruby style guide by using Rubocop

Devops is a Two-Way Street

- It's great when developers care about
 - **uptime!**
 - **scaling!**
 - **deployment!**
- Put them on call!
etc. etc. etc.



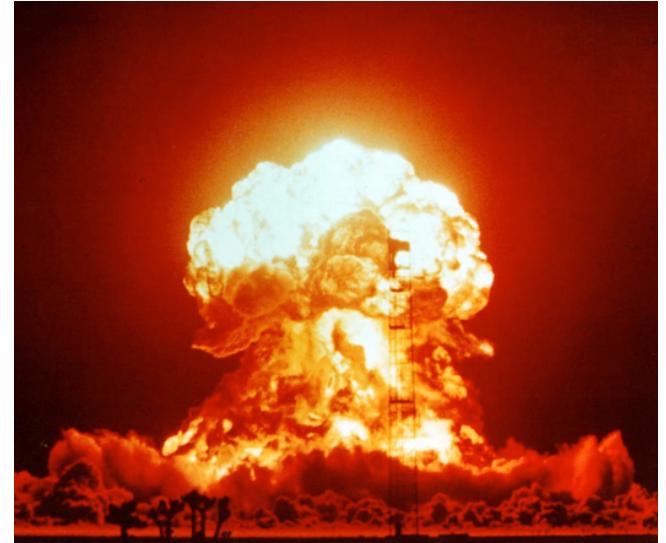
Devops is a Two-Way Street



- **Operations** also has as much or more to learn from developers as well!

Old Software Development Workflow

- Write some code
- <ad-hoc verification here>
- Go to pre-production
- <ad-hoc verification here>
- Go to production
- **Production failure**



New Software Development Workflow

- Write some code
- Write and run some **unit tests**
- Go to pre-production
- Run some **integration/acceptance tests**
- Go to production
- **Lowered chance of production failure**



Old Chef Cookbook Workflow

- Write some cookbook code
- <ad-hoc verification here>
- Go to pre-production
- <ad-hoc verification here>
- Go to production
- **Whoops, broke production**



New Chef Cookbook Workflow

- Write some cookbook code
- Check for **code correctness**
- Write and run some **unit tests**
- Go to pre-production
- Run some **integration tests**
- Go to production



Tools of the Trade

- **Code correctness:** Foodcritic, Rubocop
- **Unit tests:** ChefSpec
- **Integration tests** (not covered in this class): Test Kitchen, ServerSpec, BATS

Foodcritic

v1.0.0

What is Foodcritic?

- A correctness-checking tool for cookbooks
- Community best practices for cookbook style
- Ability to write and use custom rules
- <http://foodcritic.io/>



Foodcritic

- A **code linting** tool for Chef
 - Checks **code correctness**
 - **Catches common mistakes** before they cause problems
 - Extensible & Configurable
 - Create new **rulesets**
 - Ignore existing rulesets
 - Integrates nicely into **CI pipelines**

Rules and Tags

- All rules are numbered, e.g.
 - *FC002 - "Avoid string interpolation where not required"*
 - *FC034 - "Unused template variables"*
- Rules can be categorized and tagged, e.g.
 - 'style'
 - 'portability'
 - ...

FC002: Avoid string interpolation where not required

style strings

When you declare a resource in your recipes you frequently want to reference dynamic values such as node attributes. This warning will be shown if you are unnecessarily wrapping an attribute reference in a string.

Unnecessary string interpolation

This example would match the FC002 rule because the `version` attribute has been unnecessarily quoted.

```
# Don't do this
package "mysql-server" do
  version "#{node['mysql']['version']}"
  action :install
end
```

Modified version

This modified example would not match the FC002 rule:

```
package "mysql-server" do
  version node['mysql']['version']
  action :install
end
```

Abide by best practices

- There are 50+ rules in Foodcritic; your cookbooks should pass them (or have a good reason for not)
- You can write your own rules
- Extra community-contributed rules:
<http://www.foodcritic.io/#extra-rules>

The Problem and the Success Criteria

- **The Problem:** We want to make sure our Apache cookbook follows best practice for Chef code.
- **Success Criteria:** We use Foodcritic to check our cookbook before committing it.

Exercise: Install Foodcritic

```
$ gem install foodcritic --no-ri --no-rdoc -v 3.0.3
```

```
Fetching: foodcritic-3.0.3.gem (100%)
Successfully installed foodcritic-3.0.3
1 gem installed
```

Exercise: Check cookbook correctness with Foodcritic

```
$ foodcritic cookbooks/apache
```

```
FC003: Check whether you are running with chef server before using server-specific
features: ./recipes/default.rb:43
FC003: Check whether you are running with chef server before using server-specific
features: ./recipes/iplogger.rb:1
FC008: Generated cookbook metadata needs updating: ./metadata.rb:2
FC008: Generated cookbook metadata needs updating: ./metadata.rb:3
FC009: Resource attribute not recognised: ./recipes/default.rb:27
FC016: LWRP does not declare a default action: ./resources/vhost.rb:1
```

Exercise: Narrow the tests

```
$ foodcritic cookbooks/apache -t ~FC003 -t ~FC009
```

```
FC008: Generated cookbook metadata needs updating: ./metadata.rb:2
FC008: Generated cookbook metadata needs updating: ./metadata.rb:3
FC016: LWRP does not declare a default action: ./resources/vhost.rb:1
```

Exercise: Fix the cookbook to clear the errors

```
$ foodcritic cookbooks/apache -t ~FC003 -t ~FC009
```

```
(no output)
```

- Edit your cookbook to fix the **FC008 & FC016** errors, then rerun your Foodcritic
- Refer to <http://www.foodcritic.io/> to find the errors
- No output means no errors!

Running certain tests

- Run only certain categories of test

```
$ foodcritic <path> -t style
```

- Run specific tests

```
$ foodcritic <path> -t FC034
```

- Exclude specific tests

```
$ foodcritic <path> -t ~FC034
```

Best Practice: Run Foodcritic Before Each Commit

- Always **check** in correct code!
- Make **Foodcritic** a part of your **build pipeline** with commit hooks or other methods
- The Foodcritic web page has extensive documentation & examples showing how to build Foodcritic rules into Jenkins or Travis-CI

Rubocop

v1.0.0

Rubocop

- A robust Ruby static **code analyzer**, based on the community Ruby **style guide**
- Looks at cookbooks for **Ruby** best practices, not the **Chef DSL** - that's for Foodcritic!
- Similar to Foodcritic there are already lots of rules defined: <https://github.com/bbatsov/rubocop/blob/master/config/enabled.yml>

The Problem and the Success Criteria

- **The Problem:** We want to make sure our Apache cookbook follows best practice for Ruby code.
- **Success Criteria:** We use Rubocop to check our cookbook before committing it.

Exercise: Install RuboCop

```
$ gem install rubocop --no-ri --no-rdoc -v 0.18.1
```

```
Successfully installed rubocop-0.18.1
1 gem installed
```

Exercise: Run Rubocop on Parent Directory

```
$ rubocop cookbooks
```

```
Inspecting 92 files
.....C.....C....W...C.....C....C....C....W...CC.....C.....CCCC...C.C.WC.....
Offences:

cookbooks/chef-client/attributes/default.rb:34:1: C: Favor modifier if usage when you have a
single-line body. Another good alternative is the usage of control flow &&/|||.
if Chef::Config.has_key?(:client_fork)
^
cookbooks/chef-client/attributes/default.rb:34:17: C: has_key? is deprecated in favor of key?.
if Chef::Config.has_key?(:client_fork)
      ^^^^^^^^
cookbooks/chef-client/attributes/default.rb:57:18: C: Avoid comma after the last item of a hash.
'mailto' => nil,
^
```

Exercise: Run Rubocop on Apache

```
$ rubocop cookbooks/apache
```

```
Inspecting 8 files
CWCWCCCC

Offences:

cookbooks/apache/attributes/default.rb:1:1: C: Missing utf-8 encoding comment.
default["apache"]["indexfile"] = "index1.html"
^
cookbooks/apache/attributes/default.rb:1:9: C: Prefer single-quoted strings when you don't need
string interpolation or special symbols.
default["apache"]["indexfile"] = "index1.html"
^^^^^^^
cookbooks/apache/attributes/default.rb:1:19: C: Prefer single-quoted strings when you don't need
string interpolation or special symbols.
default["apache"]["indexfile"] = "index1.html"
^^^^^^^^^^^^^
```

Understanding the Output

Inspecting 8 files
CWCWCCCC

Offences:

- Each character represents a file checked
 - A dot represents a clean file
 - A capital letter means an offense - could be **C**onvention, **W**arning, **E**rror or **F**atal
- All errors are then listed in detail beneath this summary

Create a boilerplate .rubocop.yml



OPEN IN EDITOR: cookbooks/apache/.rubocop.yml

```
AlignParameters:  
  Enabled: false
```

```
Encoding:  
  Enabled: false
```

```
LineLength:  
  Max: 200
```

```
StringLiterals:  
  Enabled: false
```

- Turn off common, annoying warnings

SAVE FILE !

311

Exercise: Run Rubocop

```
$ rubocop cookbooks/apache
```

```
Inspecting 8 files
.C.....  
  
Offences:  
  
cookbooks/apache/files/default/plugins/modules.rb:11:48: C: Space between { and |
missing.
    fullkey, value = line.split(/\(/, 2).map { |i| i.strip}
                           ^^^  
cookbooks/apache/files/default/plugins/modules.rb:11:60: C: Space missing inside }.
    fullkey, value = line.split(/\(/, 2).map { |i| i.strip}
                           ^  
  
8 files inspected, 2 offences detected
```

Exercise: Make corrections in offending file



OPEN IN EDITOR: apache/files/default/plugins/modules.rb

```
...
modules.stdout.each_line do |line|
  fullkey, value = line.split(/\n/, 2).map { |i| i.strip }
  apache_mod = fullkey.gsub(/_module/, "")
```

SAVE FILE !

313

Exercise: Run Rubocop

```
$ rubocop cookbooks/apache
```

```
Inspecting 8 files
```

```
.....
```

```
8 files inspected, no offences detected
```

Review Questions

- Why should you check for correctness in your cookbooks?
- What is the difference between Foodcritic and Rubocop?
- At what point in the development lifecycle would you use Foodcritic and Rubocop?

An Introduction to ChefSpec

Unit testing your cookbooks to prevent regressions

v1.1.4

Lesson Objectives

- After completing the lesson, you will be able to:
 - Understand what unit testing means for Chef cookbooks and recipes
 - Understand why to write unit tests for Chef recipes
 - Use ChefSpec to create and manage a test suite for your cookbooks

Why Write Unit Tests?

- Fixing bugs before deploying code is cheap
- Fixing them afterwards is expensive
 - Programmer cost
 - Operational cost (bugs cause outages)
- Unit tests assert your intended behavior
- Unit tests run quickly
- If you need to refactor your cookbook, tests ensure you don't break anything unless you meant to

Problem Statement

- **Problem:** We broke our motd cookbook one too many times
- **Proposed Solution:** Use ChefSpec to write tests to ensure the code is valid

Install ChefSpec

```
$ gem install chefspec -v 4.0.1
```

```
Successfully installed chefspec-4.0.1  
1 gem installed
```

Introduction to ChefSpec Syntax

- ChefSpec is built on-top of RSpec
 - The standard Ruby testing tool
- RSpec has a familiar, English-like syntax, and so does ChefSpec!
 - `context` - group related tests together
 - `describe` - a test
 - `it` - a block describing some behavior
 - `expect` - expectations to assert

General Test Approach

- Set up the test
 - Make a Chef run in memory
 - Set up test harness if necessary
- Make some assertions (expectations)

General Test Format for ChefSpec

```
require 'spec_helper'

describe 'the_cookbook::default' do
  chef_run = ChefSpec::Runner.new.converge(described_recipe)
  it 'does something' do
    expect(chef_run).to some_condition
  end
end
```

- `require 'chefspec'` - load ChefSpec
- `describe` - the thing under test
- `chef_run` - create a Chef run in memory & converge it
 - `described_recipe` is syntactic sugar for '`the_cookbook::default`'
- `it` block - make some assertions (expectations)
 - `some_condition` - known as a "matcher"

Exercise: Create a spec helper



OPEN IN EDITOR: cookbooks/motd/spec/spec_helper.rb

```
require 'chefspec'  
  
at_exit { ChefSpec::Coverage.report! }
```

SAVE FILE!

- By convention, test suites have a “helper”
- Avoid restating `require 'chefspec'` over and over
- Can configure RSpec in here (formatting, enforced style, etc.)

Exercise: Make a ‘spec’ directory

```
$ mkdir cookbooks/motd/spec
```

(No output)

Exercise: Create a skeleton test



OPEN IN EDITOR: cookbooks/motd/spec/unit/default_spec.rb

```
require_relative '../spec_helper.rb'

describe 'motd::default' do
  let(:chef_run) { ChefSpec::Runner.new.converge(described_recipe) }

  it 'does something' do
    skip 'need to write this test'
  end
end
```

SAVE FILE!

- Test file name should match recipe name (`default_`), and end in `_spec.rb`
- `describe` is always the cookbook name and recipe name
- `skip` - special syntax to tell RSpec that you know you need to do some work yet

Exercise: Run rspec from the cookbook

```
$ rspec cookbooks/motd
```

```
*
```

```
Pending:
```

```
motd::default does something
# need to write this test
# ./cookbooks/motd/spec/unit/default_spec.rb:7
```

```
Finished in 0.00033 seconds
1 example, 0 failures, 1 pending
```

```
No Chef resources found, skipping coverage calculation...
```

Exercise: Write a real test



OPEN IN EDITOR: cookbooks/motd/spec/unit/default_spec.rb

```
require_relative '../spec_helper.rb'

describe 'motd::default' do
  let(:chef_run) { ChefSpec::Runner.new.converge(described_recipe) }

  it 'creates an motd correctly' do
    expect(chef_run).to create_template('/etc/motd').with(
      :user => 'root',
      :group => 'root',
      :mode => '0644'
    )
  end
end
```

SAVE FILE !

328

Exercise: Run rspec

```
$ rspec cookbooks/motd
```

Failures:

```
1) motd::default creates an motd correctly
   Failure/Error: expect(chef_run).to create_template('/etc/motd').with(
     expected "template[/etc/motd]" to have parameters:

     user "root", was nil
     group "root", was nil
   # ./spec/unit/default_spec.rb:7:in `block (2 levels) in <top (required)>'
```

Finished in 0.03019 seconds

1 example, 1 failure

...

Exercise: Fix original recipe

 **OPEN IN EDITOR:** cookbooks/motd/recipes/default.rb

```
template "/etc/motd" do
  source "motd.erb"
  mode "0644"
  owner "root"
  group "root"
end
```

- Add owner and group so the test passes.

SAVE FILE!

Exercise: Run rspec again

```
$ rspec cookbooks/motd
```

```
.
```

```
Finished in 0.02726 seconds
1 example, 0 failures
```

```
ChefSpec Coverage report generated...
```

```
Total Resources: 1
Touched Resources: 1
Touch Coverage: 100.0%
```

```
You are awesome and so is your test coverage! Have a fantastic day!
```

Exercise: Upload all recently changed cookbooks

```
$ knife cookbook upload apache motd
```

```
Uploading apache      [ 0.4.0 ]
Uploading motd        [ 0.1.0 ]
Uploaded 2 cookbooks.
```

Using Fauxhai to mock platforms

- ChefSpec is great for unit testing cross-platform cookbooks
- Let's add tests to our `mailx` cookbook that supports both Debian and RedHat variants.

Exercise: Create a spec helper



OPEN IN EDITOR: cookbooks/mailx/spec/spec_helper.rb

```
require 'chefspec'

at_exit { ChefSpec::Coverage.report! }
```

SAVE FILE!

- Same as before, only now we're in the mailx cookbook

Exercise: Write a real test



OPEN IN EDITOR: cookbooks/mailx/spec/unit/default_spec.rb

```
require_relative '../spec_helper.rb'

describe 'mailx::default' do
  context 'on Debian' do
    let(:chef_run) { ChefSpec::Runner.new({:platform => 'ubuntu', :version => '14.04'}) }
    .converge(described_recipe)

    it 'should install the correct packages' do
      expect(chef_run).to install_package 'mailutils'
    end
  end
end
```

- Set up different test contexts for different platforms

SAVE FILE !

335

Exercise: Write a real test



OPEN IN EDITOR: cookbooks/mailx/spec/unit/default_spec.rb

```
context 'on CentOS' do
  let(:chef_run) { ChefSpec::Runner.new({:platform => 'centos', :version =>
    '6.5'}).converge(described_recipe) }

  it 'should install the correct packages' do
    expect(chef_run).to install_package 'mailx'
  end
end
end
```

- Set up test context for CentOS

SAVE FILE !

336

Exercise: Run rspec

```
$ rspec cookbooks/mailx
```

```
..
```

```
Finished in 0.18828 seconds (files took 5.17 seconds to load)
2 examples, 1 failures
```

```
ChefSpec Coverage report generated...
```

```
Total Resources: 1
Touched Resources: 1
Touch Coverage: 100.0%
```

```
You are awesome and so is your test coverage! Have a fantastic day!
```

Exercise: Add cross-platform attributes



OPEN IN EDITOR: cookbooks/mailx/attributes/default.rb

```
case node['platform_family']
when "debian"
  default['mailutils']['mailx-package'] = "mailutils"
when "rhel"
  default['mailutils']['mailx-package'] = "mailx"
end
```

SAVE FILE!

Exercise: Install the package



OPEN IN EDITOR: cookbooks/mailx/recipes/default.rb

```
package node[ 'mailutils' ]['mailx-package'] do
  action :install
end
```

SAVE FILE !

Exercise: Run rspec

```
$ rspec cookbooks/mailx
```

```
..
```

```
Finished in 0.18828 seconds (files took 5.17 seconds to load)
2 examples, 0 failures
```

```
ChefSpec Coverage report generated...
```

```
Total Resources: 2
Touched Resources: 2
Touch Coverage: 100.0%
```

```
You are awesome and so is your test coverage! Have a fantastic day!
```

Review Questions

- What is ChefSpec used for?
- What tool is ChefSpec based on?
- What directory does your tests go into?
- Given a recipe named 'backup', what will the ChefSpec test filename be?

Further Resources

v1.1.4

Chef Fundamentals - Q&A Forum

- Chef Intermediate Google Group Q&A Forum
- <http://bit.ly/ChefIntermediateForum>
- Join the group and post questions

Custom Resources

- **Documentation**
 - http://docs.chef.io/chef/lwrps_custom.html
- **Examples**
 - http://docs.chef.io/lwrp_yum.html
 - <https://github.com/opscode-cookbooks/powershell>
 - <https://github.com/opscode-cookbooks/aws>

Ohai Plugins

- **Documentation**

- <http://docs.chef.io/ohai.html>

- **Examples**

- <https://github.com/opscode/ohai/tree/master/lib/ohai/plugins>
 - https://github.com/SPSCommerce/chef-etchosts/blob/master/templates/default/plugins/dhcpinfo_v7.rb

Style and Correctness

- **Documentation**

- <http://docs.chef.io/foodcritic.html>
- <http://batsov.com/rubocop/>

- **Style and Correctness Rules**

- <http://foodcritic.io/>
- <https://github.com/bbatsov/ruby-style-guide>

Unit Testing

- **Documentation**

- <http://docs.chef.io/chefspec.html>
 - <https://sethvargo.github.io/chefspec/>

- **Examples**

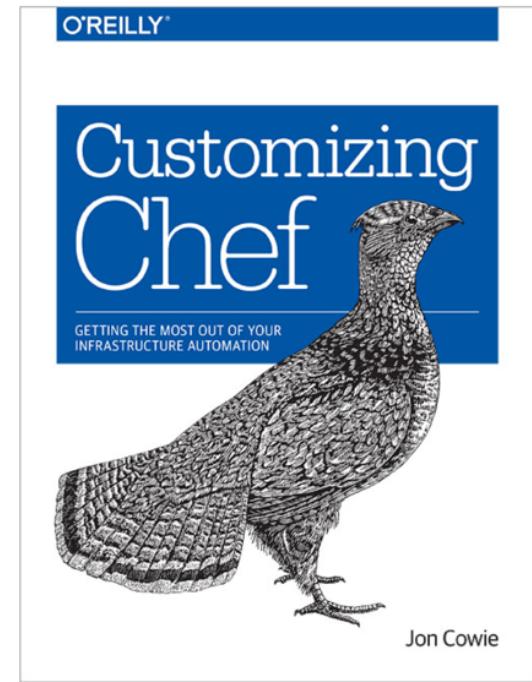
- <https://github.com/sethvargo/chefspec/tree/master/examples>

Further Resources

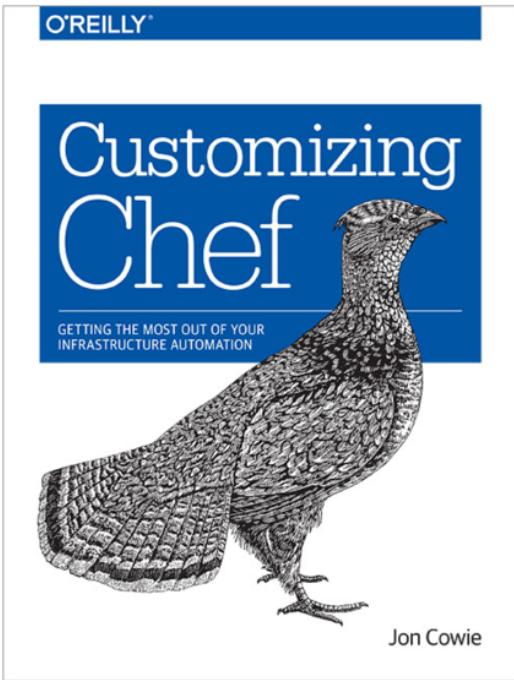
- <http://chef.io/>
- <http://supermarket.chef.io/>
- <http://docs.chef.io/>
- <http://learnchef.com>
- <http://lists.opscode.com>
- <http://youtube.com/user/getchef>
- irc.freenode.net #chef, #chef-hacking
- Twitter @chef #getchef

Learning Chef

- by Mischa Taylor & Seth Vargo of Chef
- The basics



Customizing Chef



- by Jon Cowie of Etsy
- More detailed treatment of handlers, Ohai plugins, event subsystem

Food Fight Show

- <http://foodfightshow.org>
- The Podcast Where DevOps Chef Do Battle
- Regular updates about new Cookbooks, Knife-plugins, and more
- Best Practices for working with Chef



ChefConf 2015



April 31 - March 2, 2015 | Santa Clara, CA at Convention Center

ChefConf

ChefConf 2015 is the largest, most vibrant gathering of web-scale IT and DevOps leaders, practitioners, and innovators. Featuring three days of inspired discussions, collaborative presentations, technical training, and hands-on labs focused on automating infrastructure and the continuous delivery of applications.